

LITERATURE REVIEW:

DEEP SYMBOLIC REGRESSION: RECOVERING MATHEMATICAL EXPRESSION FROM DATA VIA RISK-SEEKING POLICY GRADIENTS

NANCY LYU 02/10/2022

OUTLINE

- Motivation
- Symbolic regression
- Method
- Results.

MOTIVATION

- Understanding the underlying mathematical relationships among variables describing a dataset is a major task in the scientific process.^[1]

[1] Brenden K. Petersen, Mikel Landajuela Larra, T. Nathan Mundhenk, Claudio P. Santiago, Soo K. Kim, Joanne T. Kim, Deep symbolic regression: Recovering mathematical expressions from data via risk-seeking policy gradients.

SYMBOLIC REGRESSION

- What is symbolic regression?
- Traditional approaches to symbolic regression vs. Deep symbolic regression

WHAT IS SYMBOLIC REGRESSION?

- Symbolic regression is the process of identifying mathematical expressions that fit observed output from a black-box process. It is a discrete optimization problem generally believed to be NP-hard. [1]

[1] T. Nathan Mundhenk, Mikel Landajuela, Ruben Glatt, Claudio P. Santiago, Daniel M. Faissol, Brenden K. Petersen, Symbolic Regression via Neural-Guided Genetic Programming Population Seeding.

TRADITIONAL APPROACHES TO SYMBOLIC REGRESSION VS. DEEP SYMBOLIC REGRESSION

- The space of mathematical expressions is discrete (in model structure) and continuous (in model parameters).
- Traditional approaches to symbolic regression using evolutionary algorithms. In particular, genetic programming (GP) (Koza, 1992; Schmidt & Lipson, 2009; Back et al., 2018).
- GP exhibits high sensitivity to hyperparameters and scale poorly to larger problems.

TRADITIONAL APPROACHES TO SYMBOLIC REGRESSION VS. DEEP SYMBOLIC REGRESSION

- DSR framework: use a large model (i.e. neural network) to search the space of small models (i.e. symbolic expressions).

METHODS - Generate symbolic expression tree with RNN

- Binary tree
- Internal nodes are mathematical operators
- Terminal nodes are input variables or constants
- Pre-order traversal
- Uniqueness

Given a dataset (X, y) , $x_i \in \mathbb{R}^n$, $y_i \in \mathbb{R}$.


$$f: \mathbb{R}^n \rightarrow \mathbb{R}$$

$$\mathcal{L} = \{+, -, \times, \div, \log, \sin, \cos, x, \text{const}\}$$

$$\sin(\text{const}) / \log y$$

x_1 x_2

Expression tree

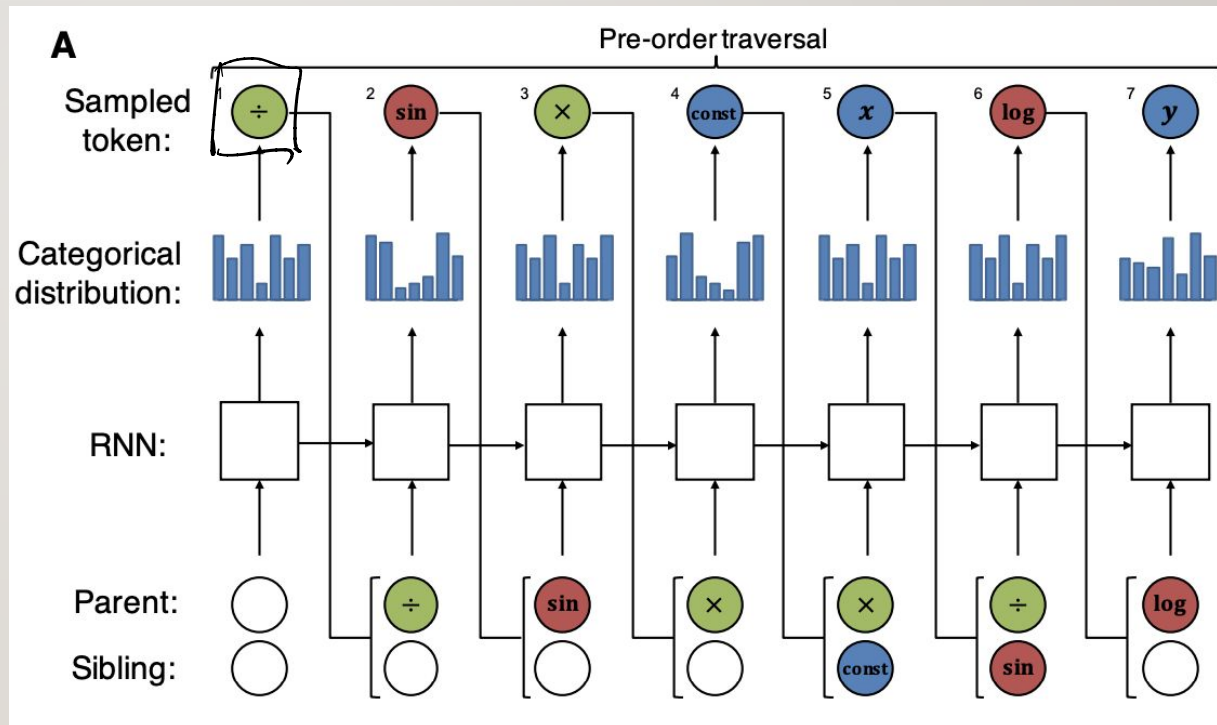


METHODS - Generate symbolic expression tree with RNN

- RNN emits a categorical distribution with parameter ψ over tokens.
- Ψ defines probabilities of selecting each token from library L .

$$p(\tau_i | \tau_{1:(i-1)}; \theta) = \psi_{\mathcal{L}(\tau_i)}^{(i)}$$


$$p(\tau | \theta) = \prod_{i=1}^{|\tau|} p(\tau_i | \tau_{1:(i-1)}; \theta) = \prod_{i=1}^{|\tau|} \psi_{\mathcal{L}(\tau_i)}^{(i)}$$




Providing hierarchical inputs to the RNN

- Parent and sibling nodes
- Empty token for node does not have a parent or sibling

Constraining the search space




- Expressions are limited to a pre-specified minimum and maximum length.
- The children of an operator should not all be constants, as the result would simply be a different constant.
- The child of a unary operator should not be the inverse of that operator, e.g. $\log(\exp(x))$ is not allowed.
- Descendants of trigonometric operators should not be trigonometric operators, e.g. $\sin(x + \cos(x))$



Reward function

- Normalized root-mean-square error (NRMSE)
- Given a dataset (X, y) of size n and candidate expression f :

$$\text{NRMSE} = \frac{1}{\sigma_y} \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - f(X_i))^2} \simeq 0$$

$$\text{R}(\tau) = \frac{1}{1 + \text{NRMSE}}$$


METHODS - training the RNN using policy gradients

- Using reinforcement learning to train the RNN to produce better-fitting expressions.
- $p(\tau | \theta)$ is like a policy
- Sampled tokens are like actions
- Standard policy gradient vs. Risk-seeking policy gradient

Standard policy gradient

- The expectation of a reward function $R(\tau)$ under expressions from the policy, $J_{\text{std}}(\theta)$.
- The standard REINFORCE policy gradient (Williams, 1992) can be used to maximize this expectation via gradient ascent:

$$\begin{aligned}\nabla_{\theta} J_{\text{std}}(\theta) &= \nabla_{\theta} \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau)] \\ &= \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau) \nabla_{\theta} \log p(\tau|\theta)]\end{aligned}$$

- Unbiased estimate of $\nabla_{\theta} J_{\text{std}}(\theta)$:

$$\nabla_{\theta} J_{\text{std}}(\theta) \approx \frac{1}{N} \sum_{i=1}^N R(\tau^{(i)}) \nabla_{\theta} \log p(\tau^{(i)}|\theta)$$

Risk-seeking policy gradient

- Optimizing the average performance or maximizing the best-case performance?
- Define $R_\varepsilon(\theta)$ as the $(1 - \varepsilon)$ -quantile of the distribution of rewards under the current policy.

$$J_{\text{risk}}(\theta; \varepsilon) \doteq \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau) \mid R(\tau) \geq R_\varepsilon(\theta)]$$

\mathcal{E}

Risk-seeking policy gradient

- How to estimate this objective via Monte Carlo sampling?
- **Proposition 1.** Let $J_{risk}(\theta; \varepsilon)$ denote the conditional expectation of rewards above the $(1-\varepsilon)$ -quantile, as in Equation (1). Then the gradient of $J_{risk}(\theta; \varepsilon)$ is given by:

$$\nabla_{\theta} J_{risk}(\theta; \varepsilon) = \mathbb{E}_{\tau \sim p(\tau|\theta)}[(R(\tau) - R_{\varepsilon}(\theta)) \cdot \nabla_{\theta} \log p(\tau|\theta) \mid R(\tau) \geq R_{\varepsilon}(\theta)]$$

$$\tilde{R}_{\varepsilon}(\theta)$$

Risk-seeking policy gradient

- A simple Monte Carlo estimate of the gradient from a batch of N samples:

$$\nabla_{\theta} J_{\text{risk}}(\theta; \varepsilon) \approx \frac{1}{\varepsilon N} \sum_{i=1}^N \left[R(\tau^{(i)}) - \tilde{R}_{\varepsilon}(\theta) \right] \cdot \mathbf{1}_{R(\tau^{(i)}) \geq \tilde{R}_{\varepsilon}(\theta)} \nabla_{\theta} \log p(\tau^{(i)} | \theta).$$

- Two differences from the standard reinforcement MC estimate:
 1. it suggests a specific baseline, $R_{\varepsilon}(\theta)$, whereas the baseline for standard policy gradients is non-specific, chosen by the user;
 2. Only the top ε fraction of samples from each batch are used in the gradient computation.


RESULTS

Benchmark	Expression	DSR	PQT	VPG	GP	Eureqa	Wolfram
Nguyen-1	$x^3 + x^2 + x$	100%	100%	96%	100%	100%	100%
Nguyen-2	$x^4 + x^3 + x^2 + x$	100%	99%	47%	97%	100%	100%
Nguyen-3	$x^5 + x^4 + x^3 + x^2 + x$	100%	86%	4%	100%	95%	100%
Nguyen-4	$x^6 + x^5 + x^4 + x^3 + x^2 + x$	100%	93%	1%	100%	70%	100%
Nguyen-5	$\sin(x^2) \cos(x) - 1$	72%	73%	5%	45%	73%	2%
Nguyen-6	$\sin(x) + \sin(x + x^2)$	100%	98%	100%	91%	100%	1%
Nguyen-7	$\log(x + 1) + \log(x^2 + 1)$	35%	41%	3%	0%	85%	0%
Nguyen-8	\sqrt{x}	96%	21%	5%	5%	0%	71%
Nguyen-9	$\sin(x) + \sin(y^2)$	100%	100%	100%	100%	100%	–
Nguyen-10	$2 \sin(x) \cos(y)$	100%	91%	99%	76%	64%	–
Nguyen-11	x^y	100%	100%	100%	7%	100%	–
Nguyen-12	$x^4 - x^3 + \frac{1}{2}y^2 - y$	0%	0%	0%	0%	0%	–
Average		83.6%	75.2%	46.7%	60.1%	73.9%	–

RESULTS

$$R'(T) = R(\mathcal{U}) + N(0, \sigma^2)^{\text{SymPy}}$$

Recovery vs dataset + noise



RESULTS

