Neural Tangent Kernel (NTK) + Preliminary Results

Madison Cooley | Feb. 9, 2022

Outline

- Overview & motivation
- How do NNs act as their width goes to infinity?
- NTK definition
- Simple PDE modeling problem
- Do PINNs exhibit similar behaviour?

- Strong PINN definition
- Preliminary Results

Overview & Motivation

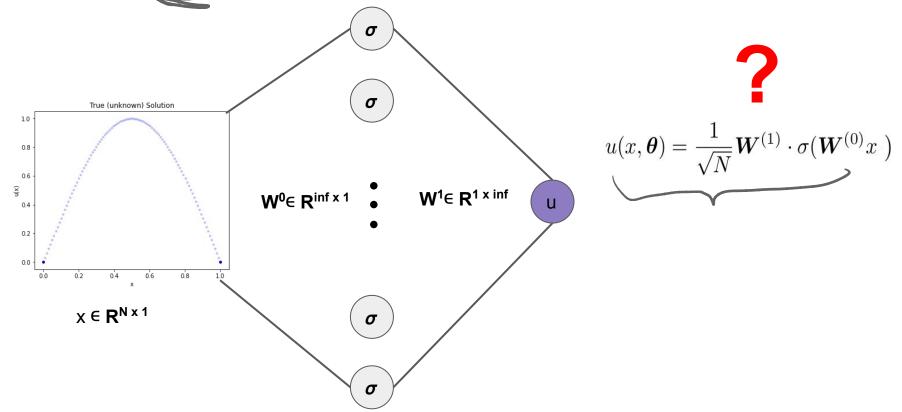
- The Neural Tangent Kernel (NTK) was proposed in 2018 by Jacot et al. for standard neural networks (NNs) [1]
- NTK provides a lens to evaluate the training performance/evolution of NNs
 - How fast is convergence? Which components converge quickest? [3]
- Many NNs and PINNs fail to provide accurate solutions on certain problems
 - Understanding why this occurs can facilitate better mitigation strategies
- Researchers propose methods that empirically fix the problem
 - Often lack theoretical justification
 - Don't address the root cause of the problem [2]
- NTK provides a way to develop methods that directly fix (some of) these issues
 - E.g. Sifan Wang, Xinling Yu, Paris Perdikaris, When and why PINNs fail to train: A neural tangent kernel perspective, 2020

^[1] Arthur Jacot, Franck Gabriel, and Clément Hongler. Neural tangent kernel: Convergence and generalization in neural networks. In Advances in neural information processing systems, 2018.

^[2] Sifan Wang, Yujun Teng, and Paris Perdikaris. Understanding and mitigating gradient pathologies in physics-informed neural networks. arXiv preprint arXiv:2001.04536, 2020.

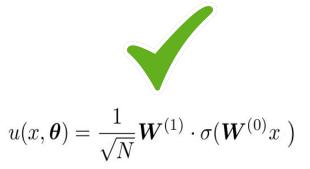
^[3] Sifan Wang, Hanwen Wang, and Paris Perdikaris. On the eigenvector bias of Fourier feature networks: From regression to solving multi-scale PDEs with physics-informed neural networks.

How do NNs act as their width goes to infinity?



How do NNs act as their width goes to infinity?

- (Jacot et al. 2018)
 - NNs are equivalent to Gaussian processes in the infinite width limit (at initialization)
 - And training evolution can be described by the NTK





Each hidden layer converges to an i.i.d centered Gaussian process [1]

Defined as the nn parameters are changing during training (for any time t)

$$Ker_t(\boldsymbol{x}, \boldsymbol{x}') = \left\langle \frac{\partial f(\boldsymbol{x}, \boldsymbol{\theta}(t))}{\partial \boldsymbol{\theta}}, \frac{\partial f(\boldsymbol{x}', \boldsymbol{\theta}(t))}{\partial \boldsymbol{\theta}} \right\rangle,$$

Defined as the nn parameters are changing during training (for any time t)

$$Ker_t(m{x},m{x}') = \left\langle \frac{\partial f(m{x},m{ heta}(t))}{\partial m{ heta}}, \frac{\partial f(m{x}',m{ heta}(t))}{\partial m{ heta}} \right
angle,$$
 erges to a deterministic kernel indization

Property 1: Kernel converges to a deterministic kernel

- At random initialization
- As width goes to infinity

Defined as the nn parameters are changing during training (for any time t)

$$Ker_t(\boldsymbol{x}, \boldsymbol{x}') = \left\langle \frac{\partial f(\boldsymbol{x}, \boldsymbol{\theta}(t))}{\partial \boldsymbol{\theta}}, \frac{\partial f(\boldsymbol{x}', \boldsymbol{\theta}(t))}{\partial \boldsymbol{\theta}} \right\rangle,$$

Property 1: Kernel converges to a deterministic kernel

- At random initialization
- As width goes to infinity

$$\begin{split} &\lim_{d_1 \to \infty} Ker_0(\boldsymbol{x}, \boldsymbol{x}') = \\ &\lim_{d_1 \to \infty} \left\langle \frac{\partial f(\boldsymbol{x}, \boldsymbol{\theta}(0))}{\partial \boldsymbol{\theta}}, \frac{\partial f(\boldsymbol{x}', \boldsymbol{\theta}(0))}{\partial \boldsymbol{\theta}} \right\rangle = \Theta^{(L)}(\boldsymbol{x}, \boldsymbol{x}') \end{split}$$

Defined as the nn parameters are changing during training (for any time t)

$$Ker_t(\boldsymbol{x}, \boldsymbol{x}') = \left\langle \frac{\partial f(\boldsymbol{x}, \boldsymbol{\theta}(t))}{\partial \boldsymbol{\theta}}, \frac{\partial f(\boldsymbol{x}', \boldsymbol{\theta}(t))}{\partial \boldsymbol{\theta}} \right\rangle,$$

Property 1: Kernel converges to a deterministic kernel

- At random initialization
- As width goes to infinity

$$\begin{split} &\lim_{d_1 \to \infty} Ker_0(\boldsymbol{x}, \boldsymbol{x}') = \\ &\lim_{d_1 \to \infty} \left\langle \frac{\partial f(\boldsymbol{x}, \boldsymbol{\theta}(0))}{\partial \boldsymbol{\theta}}, \frac{\partial f(\boldsymbol{x}', \boldsymbol{\theta}(0))}{\partial \boldsymbol{\theta}} \right\rangle = \Theta^{(L)}(\boldsymbol{x}, \boldsymbol{x}') \end{split}$$

Defined as the nn parameters are changing during training (for any time t)

$$Ker_t(oldsymbol{x},oldsymbol{x}') = \left\langle rac{\partial f(oldsymbol{x},oldsymbol{ heta}(t))}{\partial oldsymbol{ heta}}, rac{\partial f(oldsymbol{x}',oldsymbol{ heta}(t))}{\partial oldsymbol{ heta}}
ight
angle,$$

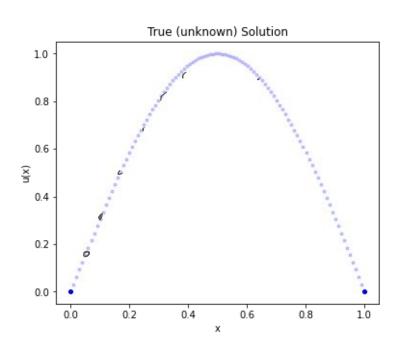
Property 1: Kernel converges to a deterministic kernel

- At random initialization
- As width goes to infinity

Property 2: Remains constant during training

$$\begin{split} &\lim_{d_1 \to \infty} Ker_0(\boldsymbol{x}, \boldsymbol{x}') = \\ &\lim_{d_1 \to \infty} \left\langle \frac{\partial f(\boldsymbol{x}, \boldsymbol{\theta}(0))}{\partial \boldsymbol{\theta}}, \frac{\partial f(\boldsymbol{x}', \boldsymbol{\theta}(0))}{\partial \boldsymbol{\theta}} \right\rangle = \Theta^{(L)}(\boldsymbol{x}, \boldsymbol{x}') \end{split}$$

1D Poisson w. Fabricated solution $u(x) = \sin(\pi x)$



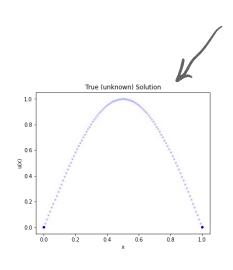
$$u_{xx}(x) = f(x), \quad \forall x \in \Omega$$

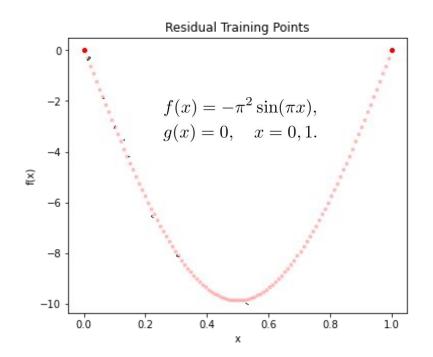
 $u(x) = g(x), \quad x \in \partial \Omega.$

$$f(x) = -\pi^2 \sin(\pi x), \quad x \in [0, 1]$$

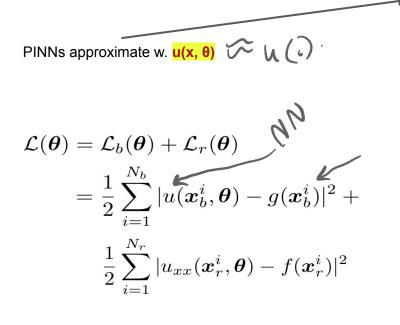
 $g(x) = 0, \quad x = 0, 1.$

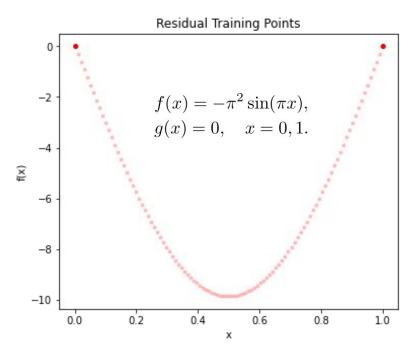
1D Poisson w. Fabricated solution $u(x) = \sin(\pi x)$



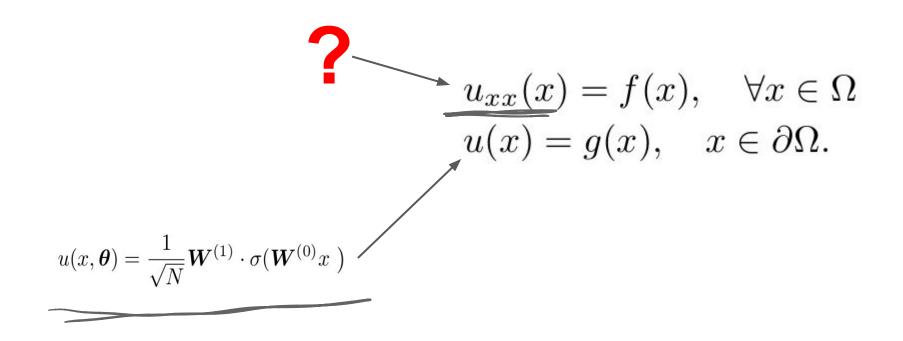


1D Poisson w. Fabricated solution $\mathbf{u}(\mathbf{x}) = \sin(\pi \mathbf{x})$





1D Poisson w. Fabricated solution $u(x) = \sin(\pi x)$



1D Poisson w. Fabricated solution $u(x) = \sin(\pi x)$

$$u_{xx}(x, oldsymbol{ heta}) = \frac{1}{\sqrt{N}} oldsymbol{W}^{(1)} \cdot \left[\ddot{\sigma}(oldsymbol{W}^{(0)}x) \odot oldsymbol{W}^{(0)} \odot oldsymbol{W}^{(0)}
ight], \qquad u_{xx}(x) = f(x), \quad orall x \in \Omega \ u(x) = g(x), \quad x \in \partial \Omega.$$

PINNs



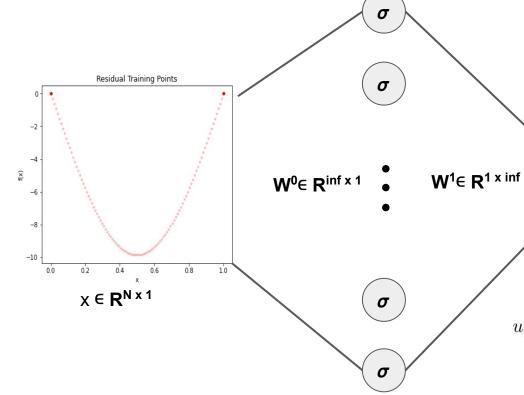
$$u(x, \boldsymbol{\theta}) = \frac{1}{\sqrt{N}} \boldsymbol{W}^{(1)} \cdot \sigma(\boldsymbol{W}^{(0)} x)$$

Property 1: NTK converges to a deterministic kernel

- At random initialization
- As width goes to infinity

u

Property 2: NTK remains constant during training



$$u_{xx}(x, \boldsymbol{\theta}) = \frac{1}{\sqrt{N}} \boldsymbol{W}^{(1)} \cdot \left[\ddot{\sigma}(\boldsymbol{W}^{(0)}x) \odot \boldsymbol{W}^{(0)} \odot \boldsymbol{W}^{(0)} \right]$$

Do PINNs exhibit similar behaviour?

- An infinitely wide NN is a Gaussian process [Jacot et al. 2020]
 - Q1: Are (infinitely wide) PINNs?
- How does the PDE residual behave in the infinite width limit
 - Q2: Converges to a deterministic kernel?
 - Q3: Remains constant during training?

WHEN AND WHY PINNS FAIL TO TRAIN: A NEURAL TANGENT KERNEL PERSPECTIVE

A PREPRINT

Sifan Wang

Graduate Group in Applied Mathematics and Computational Science University of Pennsylvania Philadelphia, PA 19104 sifanw@sas.upenn.edu

Xinling Yu

Graduate Group in Applied Mathematics and Computational Science University of Pennsylvania Philadelphia, PA 19104 xlyu@sas.upenn.edu

Paris Perdikaris

Department of Mechanichal Engineering and Applied Mechanics University of Pennsylvania Philadelphia, PA 19104 pgp@seas.upenn.edu

Do PINNs exhibit similar behaviour

- An infinitely wide NN is a Gaussian process [Jacot et al. 2020]
 - Q1: Are (infinitely wide) PINNs?
- How does the PDE residual behave in the infinite width limit
 - Q2: Converges to a deterministic kernel?
 - Q3: Remains constant during training?



Yes

WHEN AND WHY PINNS FAIL TO TRAIN: A NEURAL TANGENT KERNEL PERSPECTIVE

A PREPRINT

Sifan Wang

Graduate Group in Applied Mathematics and Computational Science University of Pennsylvania Philadelphia, PA 19104 sifanw@sas.upenn.edu

Xinling Yu

Graduate Group in Applied Mathematics and Computational Science University of Pennsylvania Philadelphia, PA 19104 xlyu@sas.upenn.edu

Paris Perdikaris

Department of Mechanichal Engineering and Applied Mechanics University of Pennsylvania Philadelphia, PA 19104 pgp@seas.upenn.edu

Under gradient flow

$$\frac{d\theta}{dt} = -\nabla \mathcal{L}(\theta)$$

$$\left[\frac{du(\boldsymbol{x}_b, \boldsymbol{\theta}(t))}{dt}\right] = -\left[\frac{\boldsymbol{K}_{uu}(t) \quad \boldsymbol{K}_{ur}(t)}{\boldsymbol{K}_{ru}(t) \quad \boldsymbol{K}_{rr}(t)}\right] \cdot \begin{bmatrix}u(\boldsymbol{x}_b, \boldsymbol{\theta}(t)) - g(\boldsymbol{x}_b)\\\mathcal{L}u(\boldsymbol{x}_r, \boldsymbol{\theta}(t)) - f(\boldsymbol{x}_r)\end{bmatrix}$$

Under gradient flow

$$\frac{d\theta}{dt} = -\nabla \mathcal{L}(\theta)$$

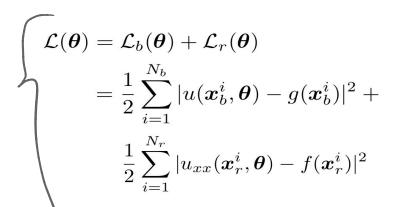
$$\Rightarrow \frac{\theta_{t+1} - \theta_t}{\eta_t} = -7 L(\theta_t)$$

$$\uparrow_{hAnixing}$$

$$\uparrow_{hAnixing}$$

Under gradient flow

$$\underbrace{\frac{d\boldsymbol{\theta}}{dt} = -\nabla \mathcal{L}(\boldsymbol{\theta})}$$



Under gradient flow

*We are mainly interested in how u(.) and uxx(.) evolve along the trajectory of gradient descent

$$\mathcal{L}(\boldsymbol{\theta}) = \mathcal{L}_b(\boldsymbol{\theta}) + \mathcal{L}_r(\boldsymbol{\theta})$$

$$= \frac{1}{2} \sum_{i=1}^{N_b} \widehat{\left(u(\boldsymbol{x}_b^i, \boldsymbol{\theta}) - g(\boldsymbol{x}_b^i)\right|^2} + \frac{1}{2} \sum_{i=1}^{N_r} \widehat{\left(u_{xx}(\boldsymbol{x}_r^i, \boldsymbol{\theta}) - f(\boldsymbol{x}_r^i)\right|^2}$$

Under gradient flow



*We are mainly interested in how u(.) and uxx(.) evolve along the trajectory of gradient descent

$$\mathcal{L}(oldsymbol{ heta}) = \mathcal{L}_b(oldsymbol{ heta}) + \mathcal{L}_r(oldsymbol{ heta})$$

$$=\frac{1}{2}\sum_{i=1}^{N_b}\widehat{u(\boldsymbol{x}_b^i,\boldsymbol{\theta})} g(\boldsymbol{x}_b^i)|^2 +$$

$$rac{1}{2}\sum_{i=1}^{N_r} |\widehat{u_{xx}}(oldsymbol{x}_r^i,oldsymbol{ heta}) - \widehat{f}(oldsymbol{x}_r^i)|^2$$

$$\frac{du(\boldsymbol{x}_b,\boldsymbol{\theta}(t))}{dt} = \frac{\gamma u}{\delta \theta} \cdot \frac{\delta \theta}{\delta t} = \frac{\gamma u}{\delta \theta} - \gamma L(\theta)$$

$$\frac{d\mathcal{L}u(\boldsymbol{x}_r,\boldsymbol{\theta}(t))}{dt}$$

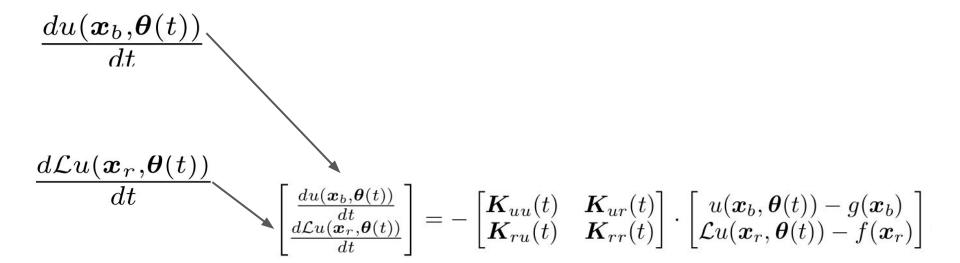
$$= -(((0)-g(0))) \cdot \left(\frac{du(0)}{d\theta}, \frac{du(0)}{d\theta}\right) = -(((0)-g(0))) \cdot \left(\frac{du(0)}{d\theta}\right)$$

Under gradient flow

*We are mainly interested in how u(.) and uxx(.) evolve along the trajectory of gradient descent

$$\mathcal{L}(\boldsymbol{\theta}) = \mathcal{L}_b(\boldsymbol{\theta}) + \mathcal{L}_r(\boldsymbol{\theta})$$

$$= \frac{1}{2} \sum_{i=1}^{N_b} |u(\boldsymbol{x}_b^i, \boldsymbol{\theta}) - g(\boldsymbol{x}_b^i)|^2 + \frac{1}{2} \sum_{i=1}^{N_r} |u_{xx}(\boldsymbol{x}_r^i, \boldsymbol{\theta}) - f(\boldsymbol{x}_r^i)|^2$$



Under gradient flow

*We are mainly interested in how u(.) and uxx(.) evolve along the trajectory of gradient descent

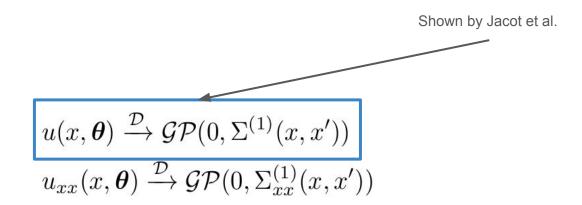
$$\mathcal{L}(\boldsymbol{\theta}) = \mathcal{L}_b(\boldsymbol{\theta}) + \mathcal{L}_r(\boldsymbol{\theta})$$

$$= \frac{1}{2} \sum_{i=1}^{N_b} |u(\boldsymbol{x}_b^i, \boldsymbol{\theta}) - g(\boldsymbol{x}_b^i)|^2 + \frac{1}{2} \sum_{i=1}^{N_r} |u_{xx}(\boldsymbol{x}_r^i, \boldsymbol{\theta}) - f(\boldsymbol{x}_r^i)|^2$$

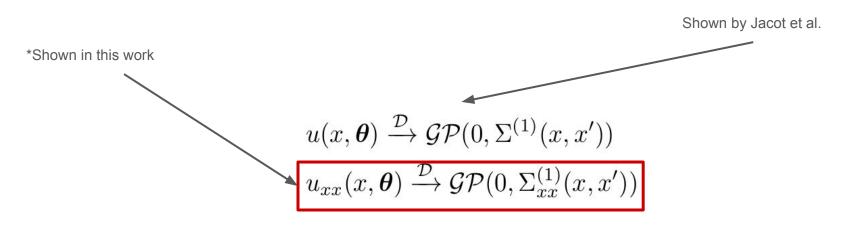
$$\frac{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt} \qquad (\boldsymbol{K}_{uu})_{ij}(t) = \left\langle \frac{du(\boldsymbol{x}_{b}^{i},\boldsymbol{\theta}(t))}{d\boldsymbol{\theta}}, \frac{du(\boldsymbol{x}_{b}^{j},\boldsymbol{\theta}(t))}{d\boldsymbol{\theta}} \right\rangle$$

$$\frac{d\mathcal{L}u(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))}{dt} \begin{bmatrix} \frac{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt} \\ \frac{d\mathcal{L}u(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))}{dt} \end{bmatrix} = -\begin{bmatrix} \boldsymbol{K}_{uu}(t) & \boldsymbol{K}_{ur}(t) \\ \boldsymbol{K}_{ru}(t) & \boldsymbol{K}_{rr}(t) \end{bmatrix} \cdot \begin{bmatrix} u(\boldsymbol{x}_{b},\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_{b}) \\ \mathcal{L}u(\boldsymbol{x}_{r},\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_{r}) \end{bmatrix}$$

Q1: Does $u_{xx}(x, \theta)$ converge in distribution to another centered Gaussian process?



Q1: Does $u_{xx}(x, \theta)$ converge in distribution to another centered Gaussian process?



- Implies model problem induces a joint GP between u (the function values) and the PDE residual at initialization
- PINNs-GP correspondence for linear PDEs

Q2: Does the NTK of PINNs converge to a deterministic kernel?

Theorem 4.3. For a physics-informed network with one hidden layer at initialization, and in the limit as the layer's width $N \to \infty$, the NTK K(t) of the PINNs model defined in equation 3.9 converges in probability to a deterministic limiting kernel, i.e,

$$\boldsymbol{K}(0) = \begin{bmatrix} \boldsymbol{K}_{uu}(0) & \boldsymbol{K}_{ur}(0) \\ \boldsymbol{K}_{ru}(0) & \boldsymbol{K}_{rr}(0) \end{bmatrix} \rightarrow \begin{bmatrix} \Theta_{uu}^{(1)} & \Theta_{ur}^{(1)} \\ \Theta_{ru}^{(1)} & \Theta_{rr}^{(1)} \end{bmatrix} := \boldsymbol{K}^*, \tag{4.7}$$



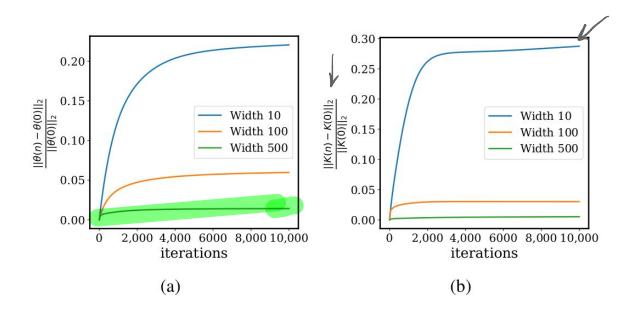
Q2: Does the NTK of PINNs converge to a deterministic kernel?

Theorem 4.3. For a physics-informed network with one hidden layer at initialization, and in the limit as the layer's width $N \to \infty$, the NTK K(t) of the PINNs model defined in equation 3.9 converges in probability to a deterministic limiting kernel, i.e,

$$\boldsymbol{K}(0) = \begin{bmatrix} \boldsymbol{K}_{uu}(0) & \boldsymbol{K}_{ur}(0) \\ \boldsymbol{K}_{ru}(0) & \boldsymbol{K}_{rr}(0) \end{bmatrix} \rightarrow \begin{bmatrix} \Theta_{uu}^{(1)} & \Theta_{ur}^{(1)} \\ \Theta_{ru}^{(1)} & \Theta_{rr}^{(1)} \end{bmatrix} := \boldsymbol{K}^*, \tag{4.7}$$

$$\lim_{N \to \infty} \sup_{t \in [0,T]} \| \mathbf{K}(t) - \mathbf{K}(0) \|_2 = 0,$$

Some numerical results



Q2: Does the NTK of PINNs converge to a deterministic kernel? Q3: Does the NTK of PINNs remain constant during training?

$$K(t) \approx K(0) \approx K^*, \quad \forall t,$$

Q2: Does the NTK of PINNs converge to a deterministic kernel? Q3: Does the NTK of PINNs remain constant during training?

$$\boldsymbol{K}(t) \approx \boldsymbol{K}(0) \approx \boldsymbol{K}^*, \quad \forall t,$$

$$\begin{bmatrix} \frac{du(\boldsymbol{x}_b,\boldsymbol{\theta}(t))}{dt} \\ \frac{d\mathcal{L}u(\boldsymbol{x}_r,\boldsymbol{\theta}(t))}{dt} \end{bmatrix} = -\begin{bmatrix} \boldsymbol{K}_{uu}(t) & \boldsymbol{K}_{ur}(t) \\ \boldsymbol{K}_{ru}(t) & \boldsymbol{K}_{rr}(t) \end{bmatrix} \cdot \begin{bmatrix} u(\boldsymbol{x}_b,\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_b) \\ \mathcal{L}u(\boldsymbol{x}_r,\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_r) \end{bmatrix}$$

Q2: Does the NTK of PINNs converge to a deterministic kernel?

$$\boldsymbol{K}(t) \approx \boldsymbol{K}(0) \approx \boldsymbol{K}^*, \quad \forall t,$$

$$\begin{bmatrix} \frac{du(\boldsymbol{x}_b,\boldsymbol{\theta}(t))}{dt} \\ \frac{d\mathcal{L}u(\boldsymbol{x}_r,\boldsymbol{\theta}(t))}{dt} \end{bmatrix} = -\begin{bmatrix} \boldsymbol{K}_{uu}(t) & \boldsymbol{K}_{ur}(t) \\ \boldsymbol{K}_{ru}(t) & \boldsymbol{K}_{rr}(t) \end{bmatrix} \cdot \begin{bmatrix} u(\boldsymbol{x}_b,\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_b) \\ \mathcal{L}u(\boldsymbol{x}_r,\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_r) \end{bmatrix}$$

$$pprox - \mathbf{K}^* \begin{bmatrix} u(\mathbf{x}_b, \boldsymbol{\theta}(t)) - g(\mathbf{x}_b) \\ u_{xx}(\mathbf{x}_r, \boldsymbol{\theta}(t)) - f(\mathbf{x}_r). \end{bmatrix}$$

Q2: Does the NTK of PINNs converge to a deterministic kernel?

$$\begin{bmatrix} \frac{du(\boldsymbol{x}_b,\boldsymbol{\theta}(t))}{dt} \\ \frac{du_{xx}(\boldsymbol{x}_r,\boldsymbol{\theta}(t))}{dt} \end{bmatrix} \approx -\boldsymbol{K}^* \begin{bmatrix} u(\boldsymbol{x}_b,\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_b) \\ u_{xx}(\boldsymbol{x}_r,\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_r). \end{bmatrix}$$

Q2: Does the NTK of PINNs converge to a deterministic kernel?

$$\begin{vmatrix} \frac{du(\boldsymbol{x}_b,\boldsymbol{\theta}(t))}{dt} \\ \frac{du_{xx}(\boldsymbol{x}_r,\boldsymbol{\theta}(t))}{dt} \end{vmatrix} \approx -\boldsymbol{K}^* \begin{bmatrix} u(\boldsymbol{x}_b,\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_b) \\ u_{xx}(\boldsymbol{x}_r,\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_r). \end{bmatrix}$$

$$\begin{bmatrix} u(\boldsymbol{x}_{test}, \boldsymbol{\theta}(\infty)) \\ u_{xx}(\boldsymbol{x}_{test}, \boldsymbol{\theta}(\infty)) \end{bmatrix} \approx \boldsymbol{K}_{test}^*(\boldsymbol{K}^*)^{-1} \cdot \begin{bmatrix} g(\boldsymbol{x}_b) \\ f(\boldsymbol{x}_r) \end{bmatrix}$$

$$\begin{bmatrix} \frac{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt} \\ \frac{dt}{du_{\boldsymbol{x}x}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))} \end{bmatrix} \approx -\boldsymbol{K}(0) \begin{bmatrix} u(\boldsymbol{x}_{b},\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_{b}) \\ u_{\boldsymbol{x}x}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_{r}) \end{bmatrix}$$

$$\boldsymbol{Q} \left(\begin{bmatrix} \frac{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt} \\ \frac{dt}{du_{\boldsymbol{\xi},\boldsymbol{x}}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))} \\ \frac{dt}{dt} \end{bmatrix} - \begin{bmatrix} g(\boldsymbol{x}_{b}) \\ f(\boldsymbol{x}_{r}) \end{bmatrix} \right) \approx -e^{-\mathbf{\Lambda}t} \boldsymbol{Q} \cdot \begin{bmatrix} g(\boldsymbol{x}_{b}) \\ f(\boldsymbol{x}_{r}) \end{bmatrix}$$

$$\begin{bmatrix} \frac{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt} \\ \frac{dt}{du_{xx}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))} \\ \frac{dt}{dt} \end{bmatrix} \approx -\boldsymbol{K}(0) \begin{bmatrix} u(\boldsymbol{x}_{b},\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_{b}) \\ u_{xx}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_{r}) \end{bmatrix} \longrightarrow \begin{bmatrix} \frac{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt} \\ \frac{du_{xx}(\frac{dt}{dx_{r},\boldsymbol{\theta}(t))}}{dt} \end{bmatrix} \approx \left(I - e^{-\boldsymbol{K}(0)t}\right) \cdot \begin{bmatrix} g(\boldsymbol{x}_{b}) \\ f(\boldsymbol{x}_{r}) \end{bmatrix}$$

$$\boldsymbol{Q} \left(\begin{bmatrix} \frac{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt} \\ \frac{du_{xx}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))}{dt} \end{bmatrix} - \begin{bmatrix} g(\boldsymbol{x}_{b}) \\ f(\boldsymbol{x}_{r}) \end{bmatrix}\right) \approx -e^{-\boldsymbol{\Lambda}t}\boldsymbol{Q} \cdot \begin{bmatrix} g(\boldsymbol{x}_{b}) \\ f(\boldsymbol{x}_{r}) \end{bmatrix}$$

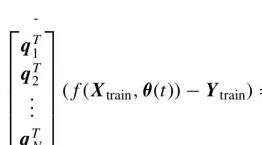
$$\begin{bmatrix} \underbrace{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}_{dt} \\ \underbrace{du_{xx}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))}_{dt} \end{bmatrix} \approx -\boldsymbol{K}(0) \begin{bmatrix} u(\boldsymbol{x}_{b},\boldsymbol{\theta}(t)) - g(\boldsymbol{x}_{b}) \\ u_{xx}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t)) - f(\boldsymbol{x}_{r}) \end{bmatrix} \longrightarrow \begin{bmatrix} \underbrace{du(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}_{dt} \\ \underbrace{du_{xx}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))}_{dt} \end{bmatrix} \approx \left(I - e^{-\boldsymbol{K}(0)t}\right) \cdot \begin{bmatrix} g(\boldsymbol{x}_{b}) \\ f(\boldsymbol{x}_{r}) \end{bmatrix}$$

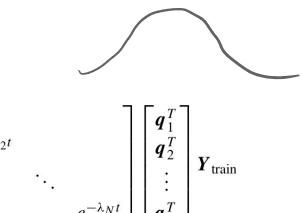
$$\begin{bmatrix} \frac{du(\boldsymbol{x}, \boldsymbol{\theta}(t))}{dt} \\ \frac{du_{\boldsymbol{x}\boldsymbol{x}}(\boldsymbol{x}_r, \boldsymbol{\theta}(t))}{dt} \end{bmatrix} - \begin{bmatrix} g(\boldsymbol{x}_b) \\ f(\boldsymbol{x}_r) \end{bmatrix} \approx \left(I - e^{-\boldsymbol{K}(0)t}\right) \cdot \begin{bmatrix} g(\boldsymbol{x}_b) \\ f(\boldsymbol{x}_r) \end{bmatrix} - \begin{bmatrix} g(\boldsymbol{x}_b) \\ f(\boldsymbol{x}_r) \end{bmatrix}$$

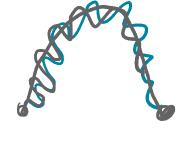
$$\approx -\boldsymbol{Q}^T e^{-\boldsymbol{\Lambda}t} \boldsymbol{Q} \cdot \begin{bmatrix} g(\boldsymbol{x}_b) \\ f(\boldsymbol{x}_r) \end{bmatrix},$$

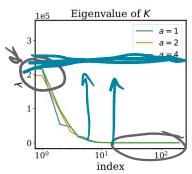


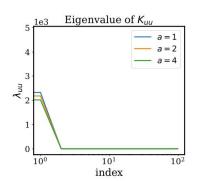
$$Q\left(\begin{bmatrix}\frac{d\nu(\boldsymbol{x}_{b},\boldsymbol{\theta}(t))}{dt}\\\frac{du_{\boldsymbol{x}\boldsymbol{x}}(\boldsymbol{x}_{r},\boldsymbol{\theta}(t))}{dt}\end{bmatrix} - \begin{bmatrix}g(\boldsymbol{x}_{b})\\f(\boldsymbol{x}_{r})\end{bmatrix}\right) \approx -e^{-\Lambda t}Q \cdot \begin{bmatrix}g(\boldsymbol{x}_{b})\\f(\boldsymbol{x}_{r})\end{bmatrix}$$

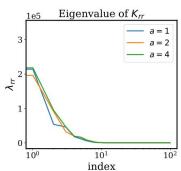














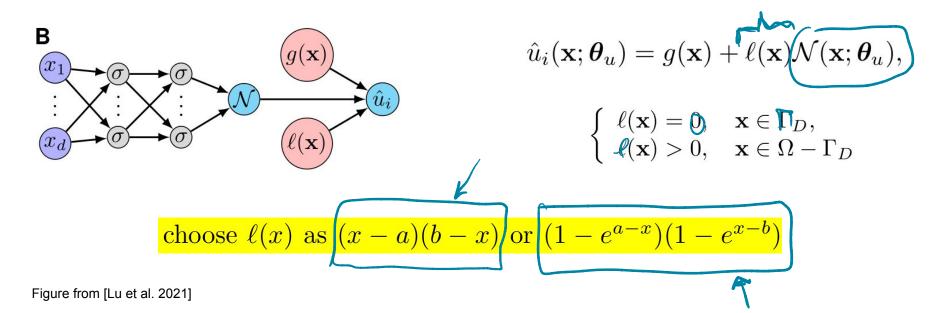
How is the NTK of PINNs practically useful?

- Define the avg. convergence rate as the mean of all eigenvalues
- Reweight loss terms given these computed averages
- Balances out loss terms such that one doesn't dominate
- Not a fool-proof solution though

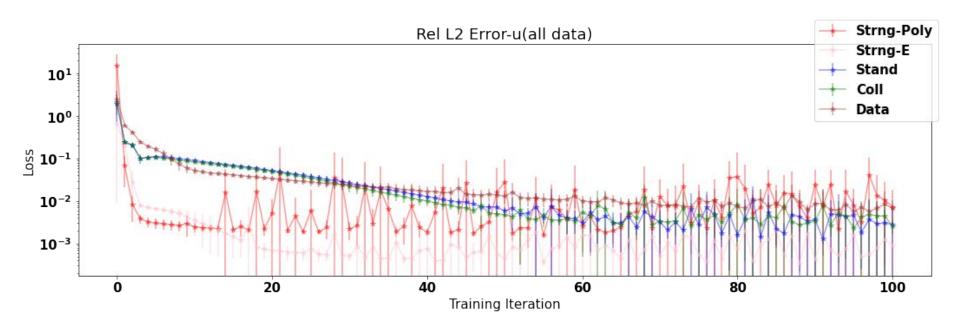
$$L = \lambda_1 f_1 + \lambda_2 f_3$$

Strong PINNs

- Enforce boundary conditions exactly by modifying architecture
- As opposed to through an additional term in the loss function



Preliminary Strong PINN Results



Preliminary Strong PINN Results

y(x) = y x y(x) = z x

