

# LLM Quantization

Xin Yu 2023.11.01



Write python codes




Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, Harsha Nori, Hamid Palangi, Marco Tulio Ribeiro, Yi Zhang

## Write LaTeX codes

**Prompt:** Draw a unicorn in TiKZ.

**GPT-4:** [Produces L<sup>A</sup>T<sub>E</sub>X compiling to following picture.]

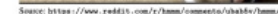


**Prompt:** Suppose  $g(x) = f^{-1}(x)$ ,  $g(0) = 5$ ,  $g(4) = 7$ ,  $g(3) = 2$ ,  $g(7) = 9$ ,  $g(9) = 6$  what is  $f(f(f(6)))$ ?

**GPT-4:** There are different ways to approach this problem, but one possible method is to use the fact that  $g(x) = f^{-1}(x)$  implies that  $f(g(x)) = x$  for any  $x$  in the domain of  $g$ . That is, applying  $f$  to both sides of the equation  $g(x) = f^{-1}(x)$  gives  $f(g(x)) = f(f^{-1}(x)) = x$ . This means that we can undo the effect of  $f$  by applying  $g$ , and vice versa. Using this idea, we can try to find  $f(6)$  by looking for the value of  $x$  such that  $g(x) = 6$ . From the given information, we see that  $g(9) = 6$ , so  $f(6) = f(g(9)) = 9$ . Similarly, we can find  $f(9)$  by looking for the value of  $x$  such that  $g(x) = 9$ . We see that  $g(7) = 9$ , so  $f(9) = f(g(7)) = 7$ . Finally, we can find  $f(7)$  by looking for the value of  $x$  such that  $g(x) = 7$ . We see that  $g(4) = 7$ , so  $f(7) = f(g(4)) = 4$ . Therefore,  $f(f(f(6))) = f(f(9)) = f(7) = 4$ .

## Explain jokes

User	What is funny about this image? Describe it panel by panel.
------	---



GPT-4 The image shows a package for a "Lightning Cable" adapter with three panels.

Panel 1: A smartphone with a VGA connector (a large, blue, 15-pin connector typically used for computer monitors) plugged into its charging port.

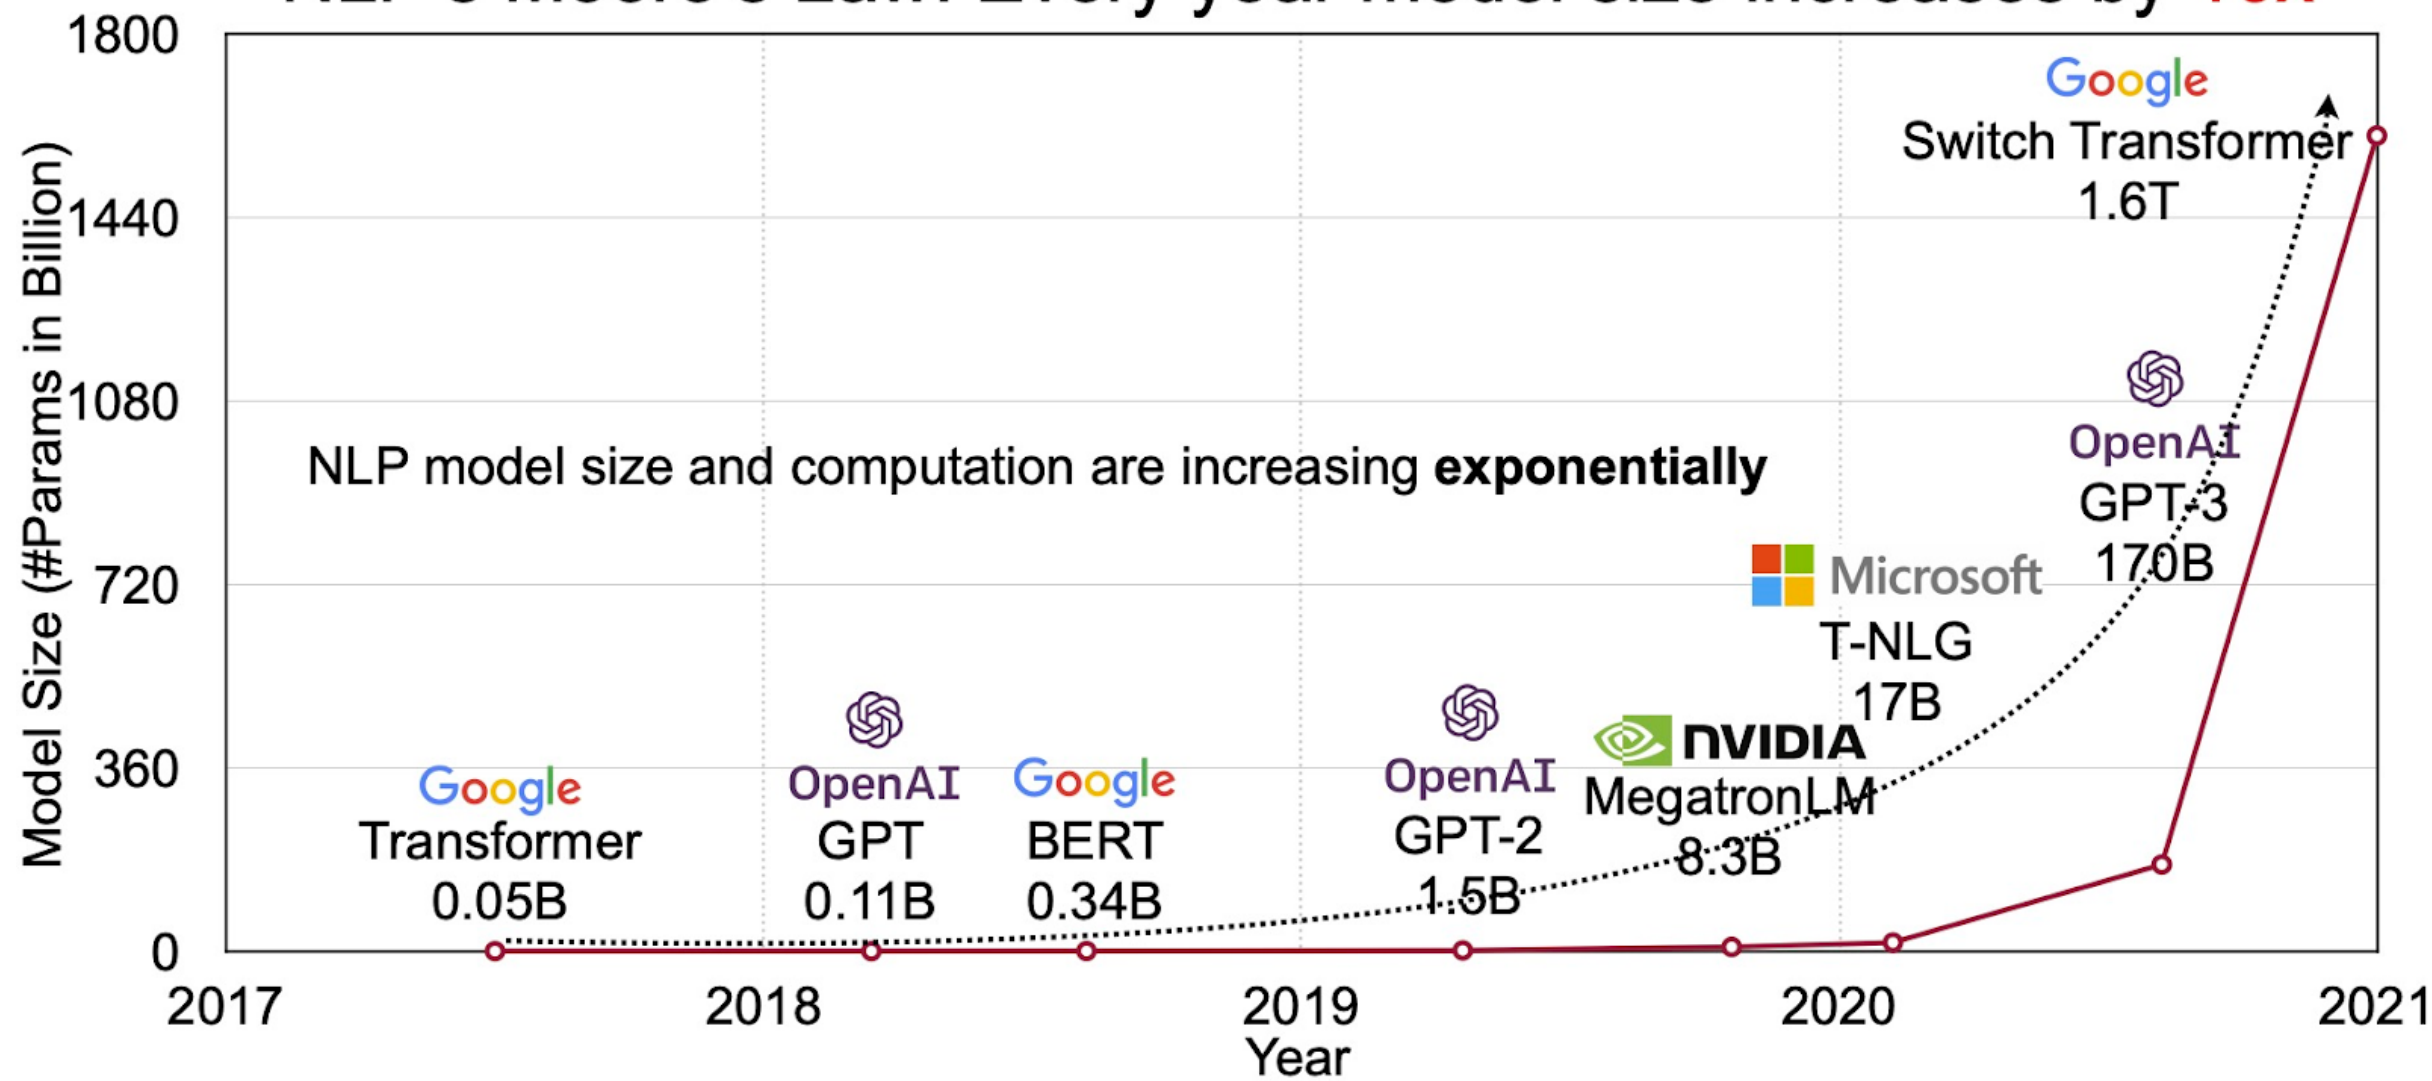
Panel 2: The package for the "Lightning Cable" adapter with a picture of a VGA connector on it.

Panel 3: A close-up of the VGA connector with a small Lightning connector (used for charging iPhones and other Apple devices) at the end.

The humor in this image comes from the absurdity of plugging a large, outdated VGA connector into a small, modern smartphone charging port.

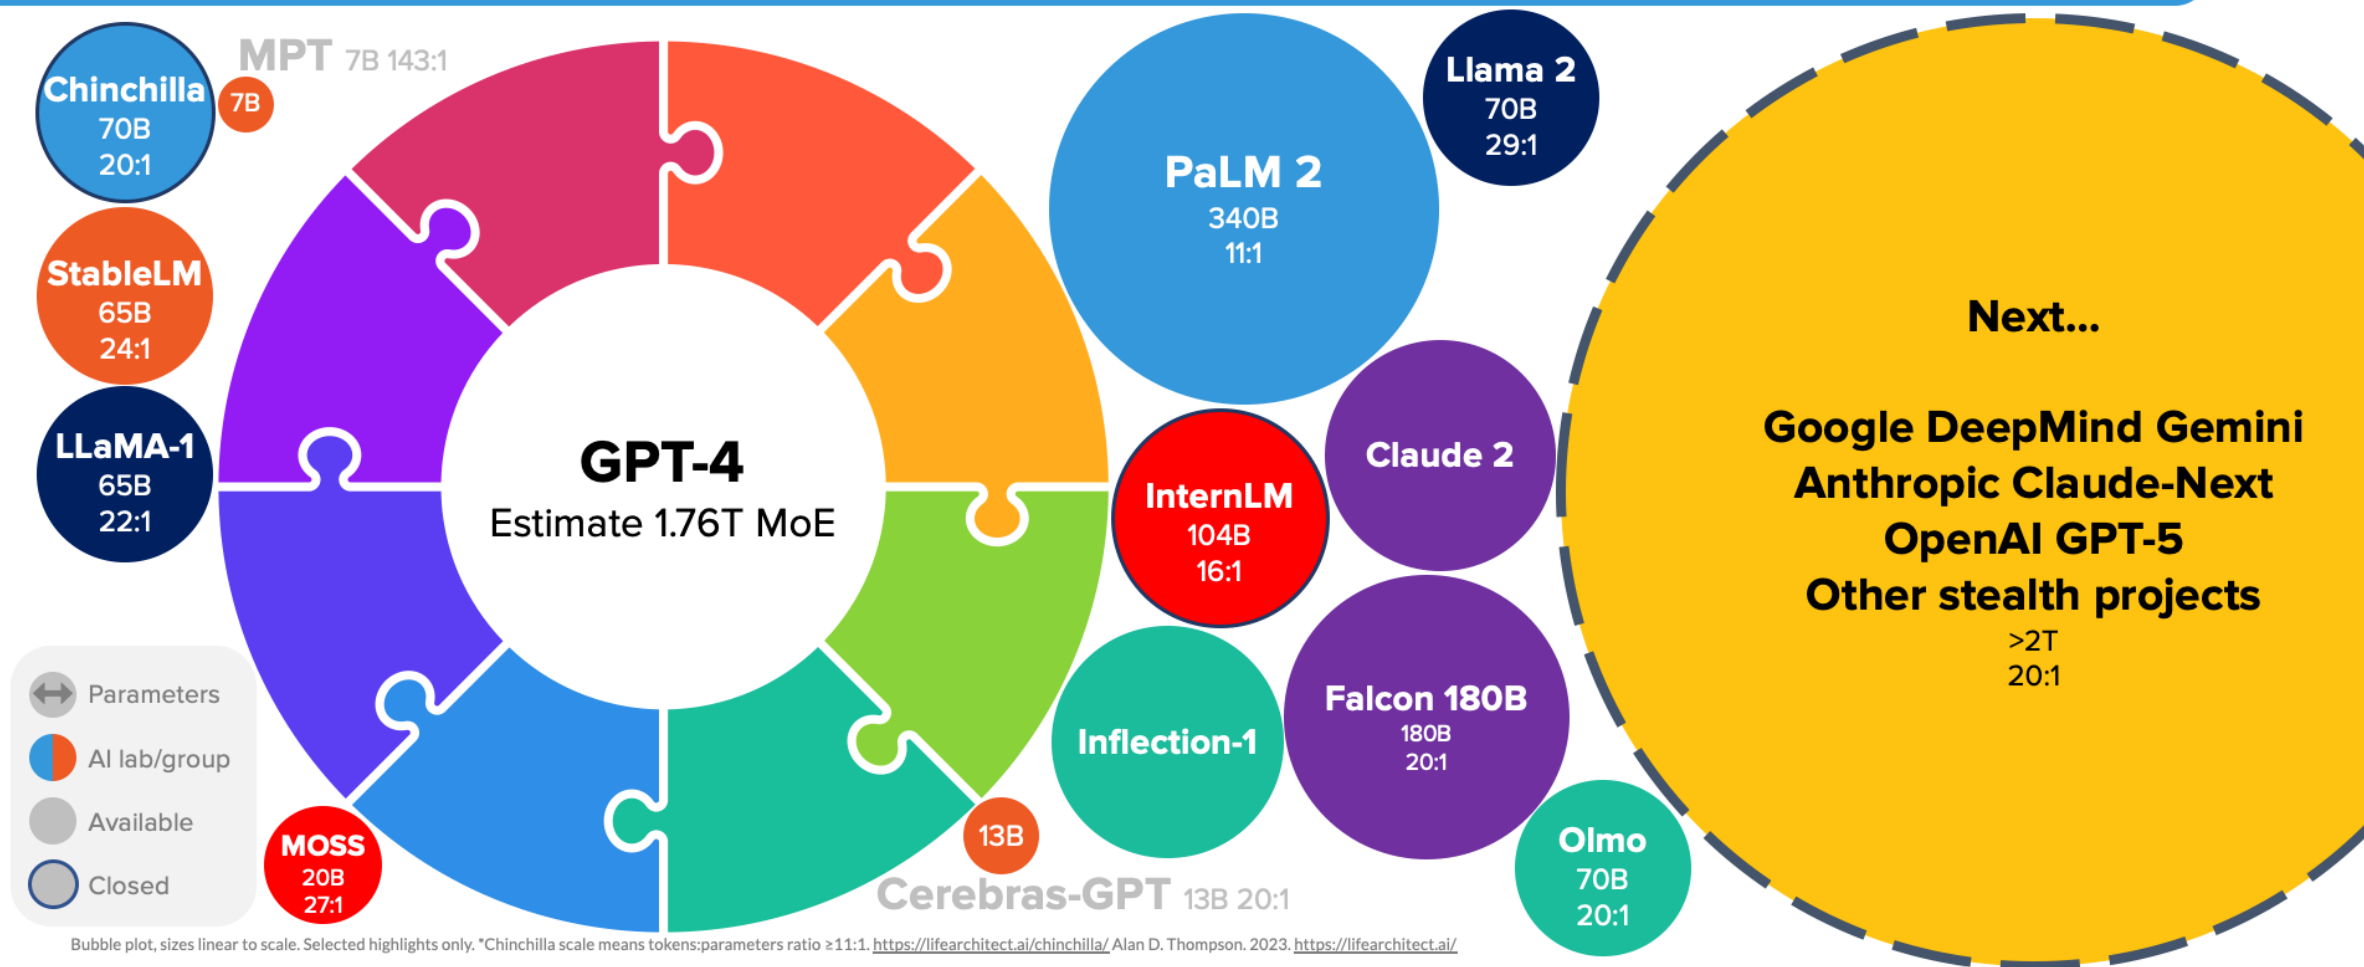
**Table 3.** Example prompt demonstrating GPT-4’s visual input capability. The prompt consists of a question about an image with multiple panels which GPT-4 is able to answer.

# NLP's Moore's Law: Every year model size increases by 10x

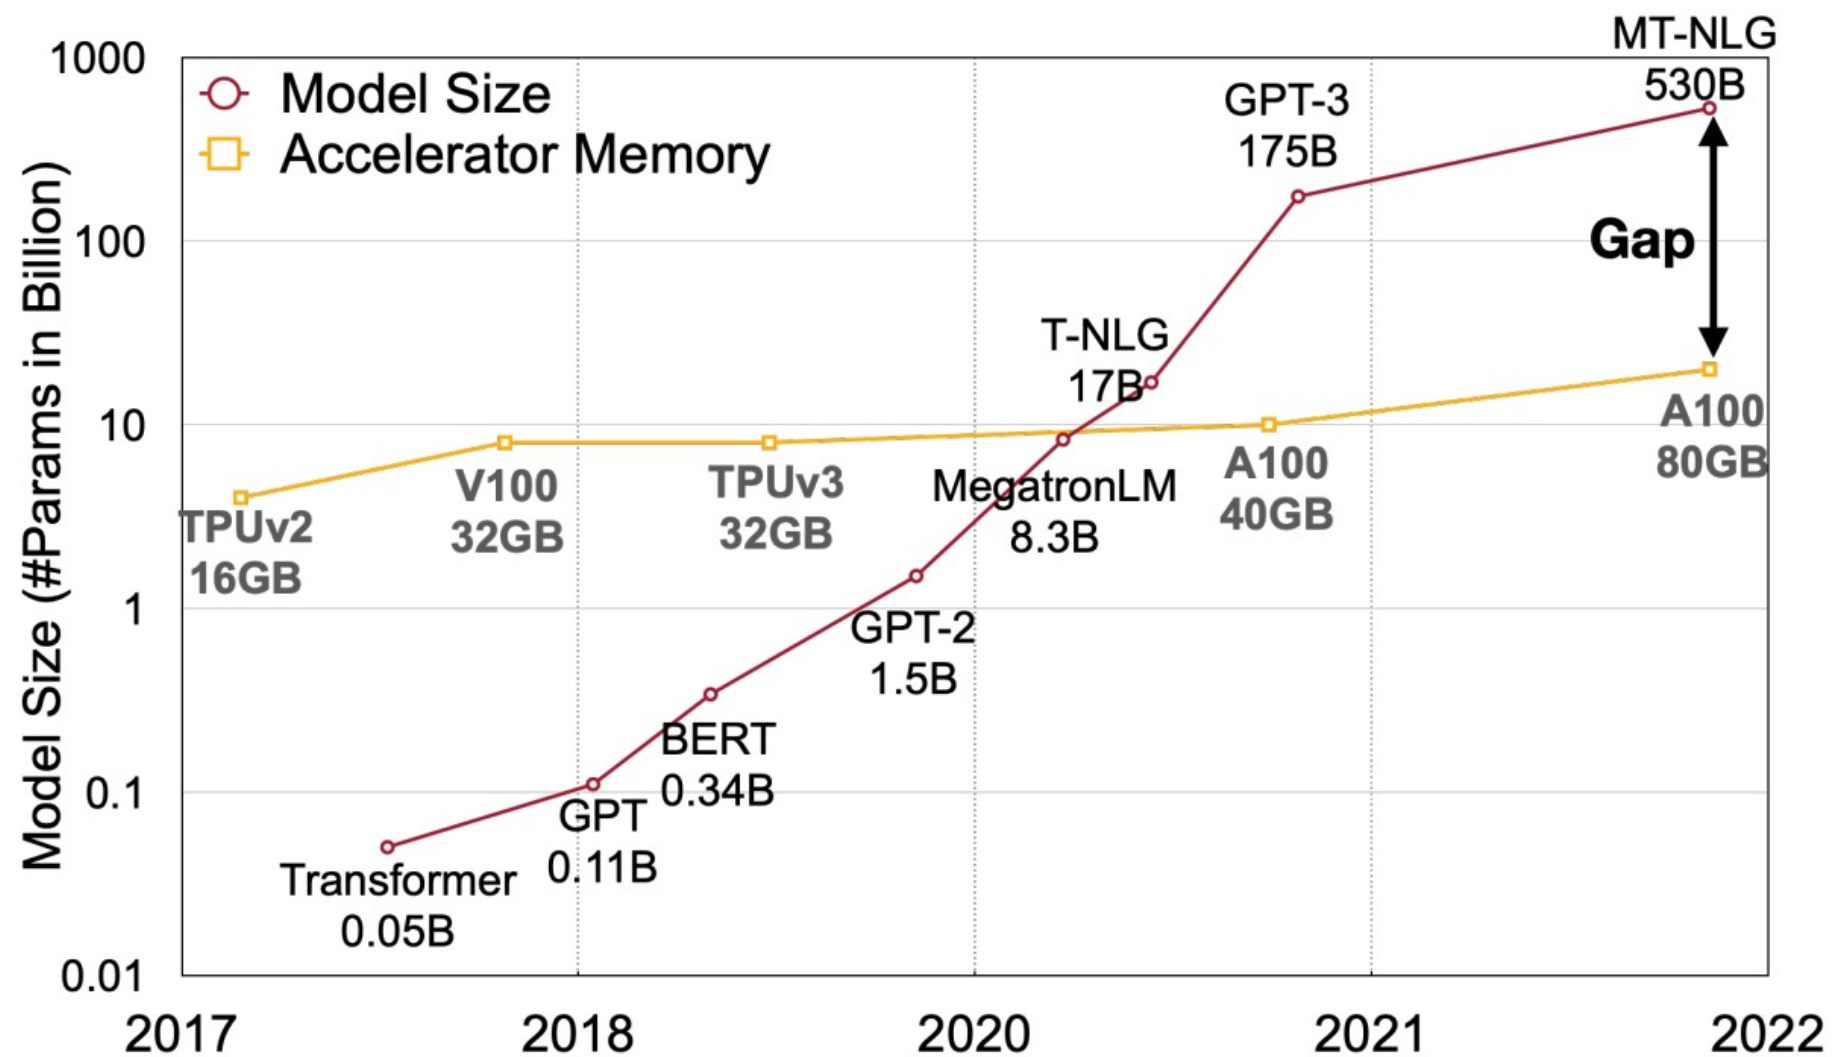


# 2023-2024 OPTIMAL LANGUAGE MODELS

SEP/  
2023



LifeArchitect.ai/models

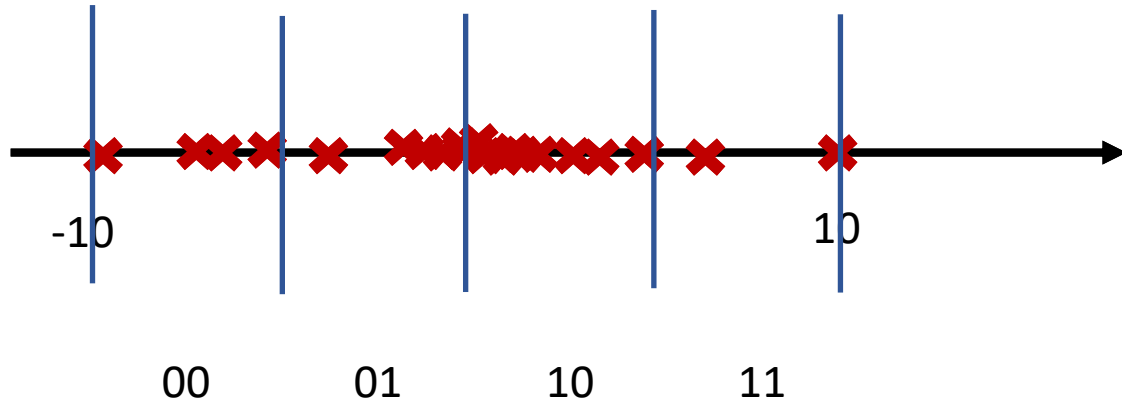


# Quantization: Data with less bits



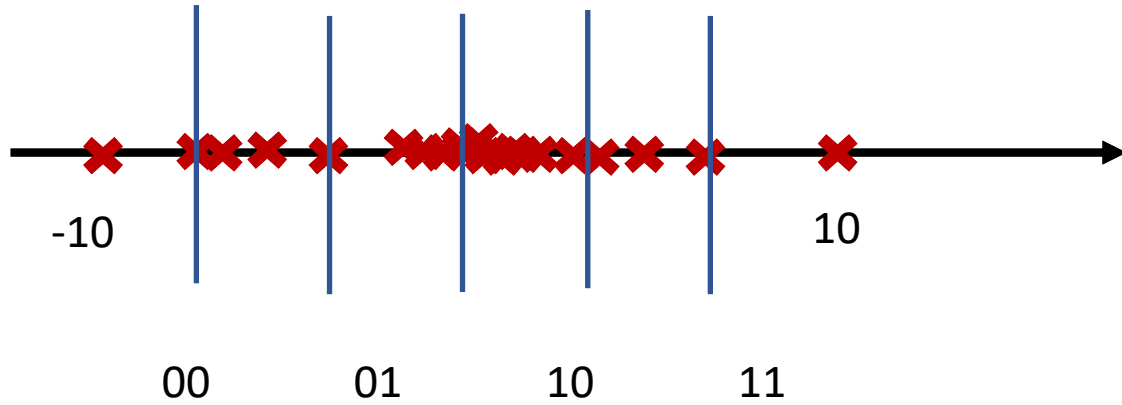
# Quantization: Data with less bits

RTN: round to nearest



# Quantization: Data with less bits

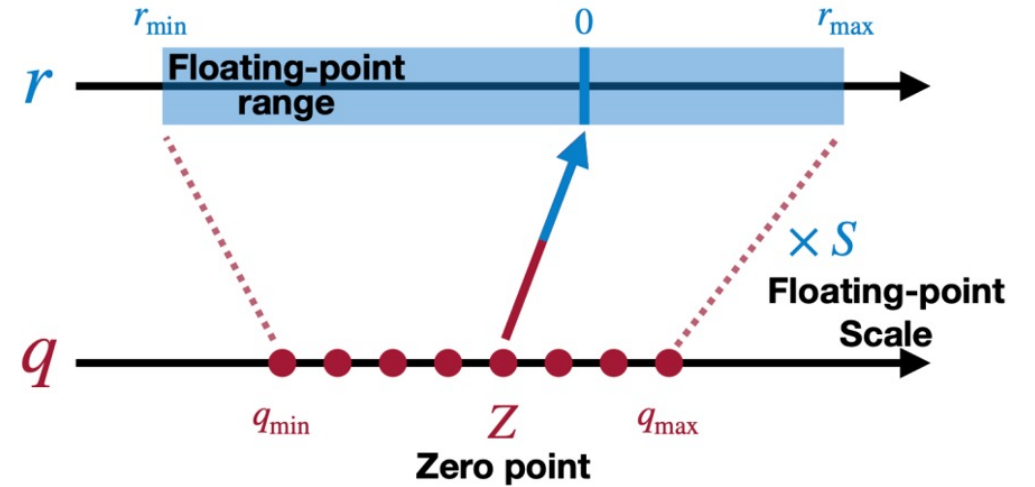
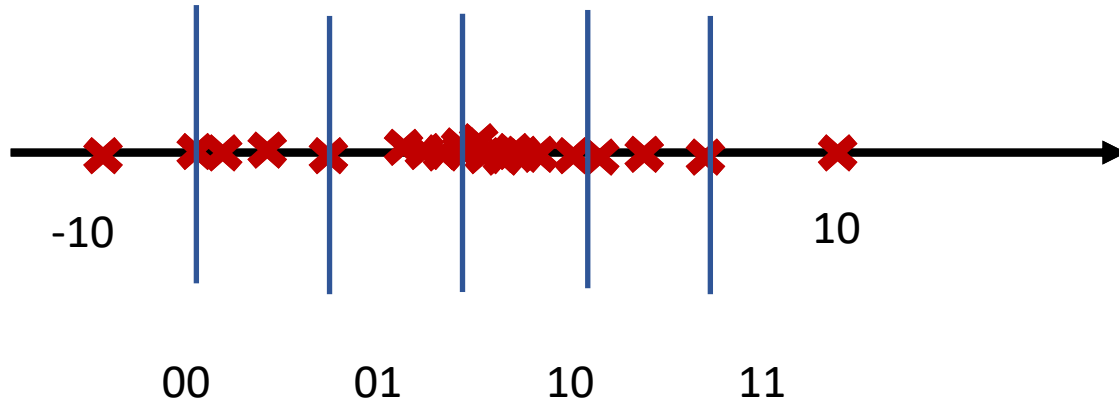
RTN: round to nearest





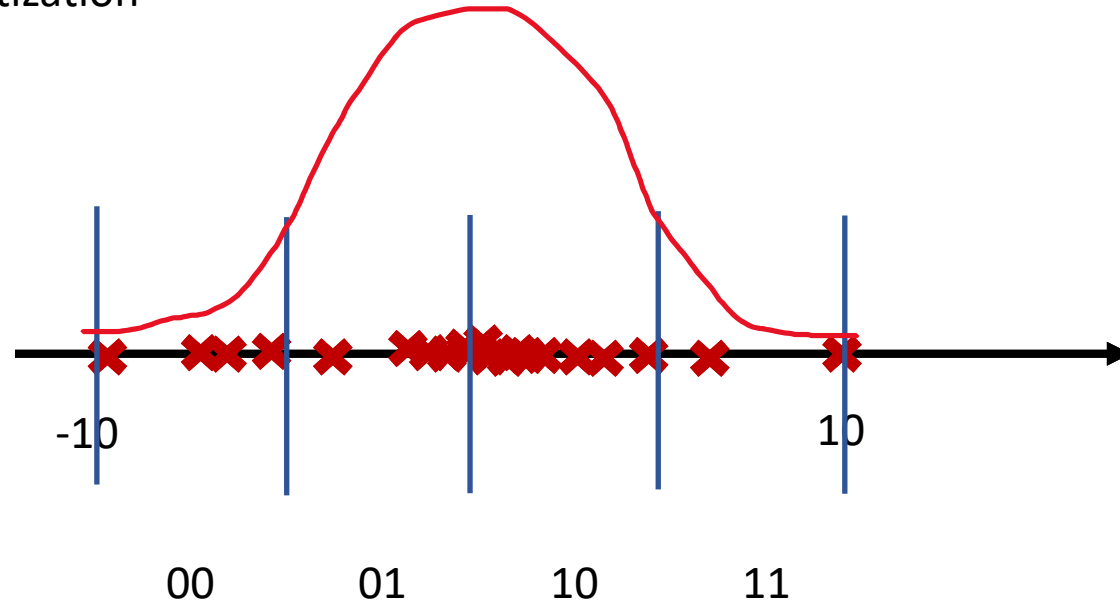
# Quantization: Data with less bits

RTN: round to nearest



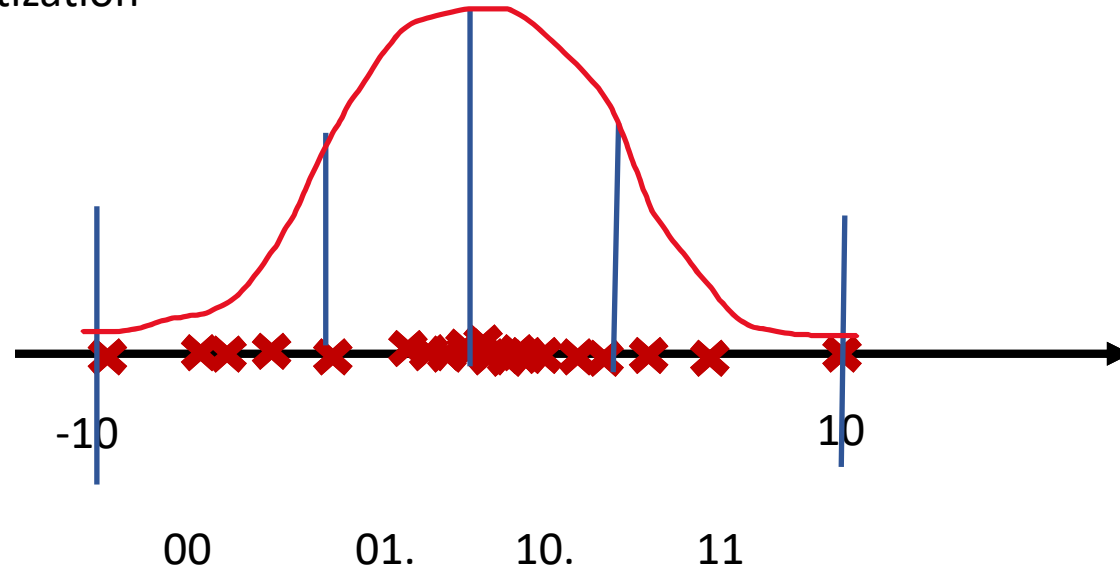
# Quantization: Data with less bits

Normal Float Quantization



# Quantization: Data with less bits

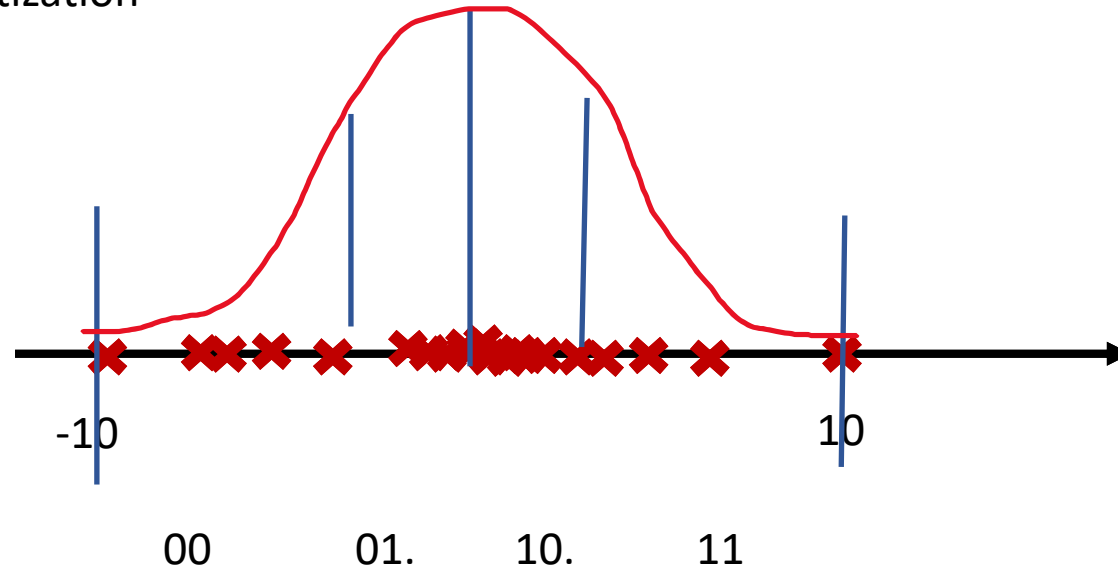
Normal Float Quantization



Quantile Quantization

# Quantization: Data with less bits

Normal Float Quantization



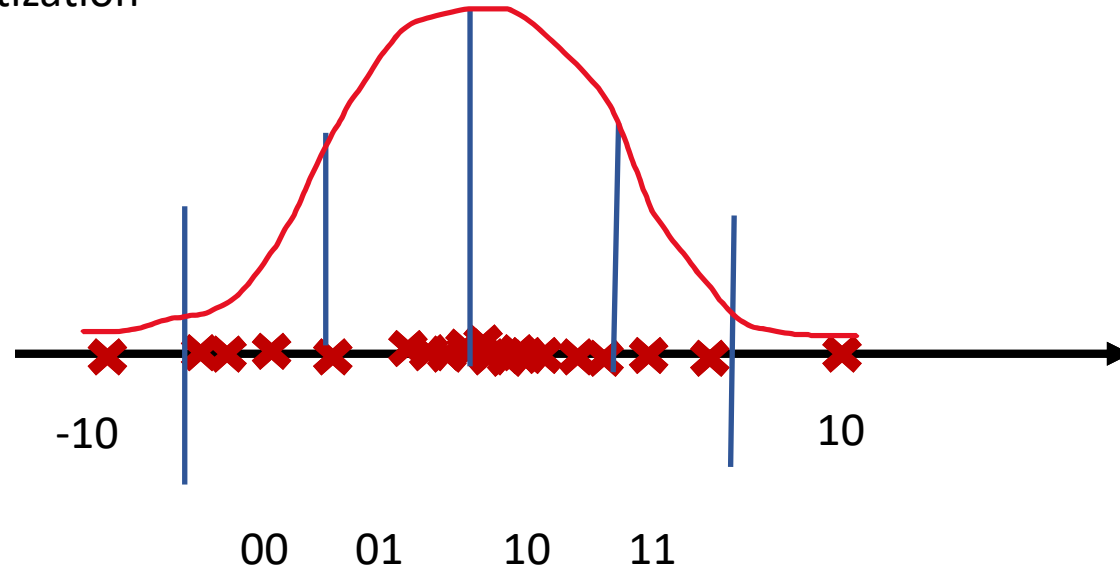
Quantile Quantization

The exact values of the NF4 data type are as follows:

[-1.0, -0.6961928009986877, -0.5250730514526367,  
-0.39491748809814453, -0.28444138169288635, -0.18477343022823334,  
-0.09105003625154495, 0.0, 0.07958029955625534, 0.16093020141124725,  
0.24611230194568634, 0.33791524171829224, 0.44070982933044434,  
0.5626170039176941, 0.7229568362236023, 1.0]

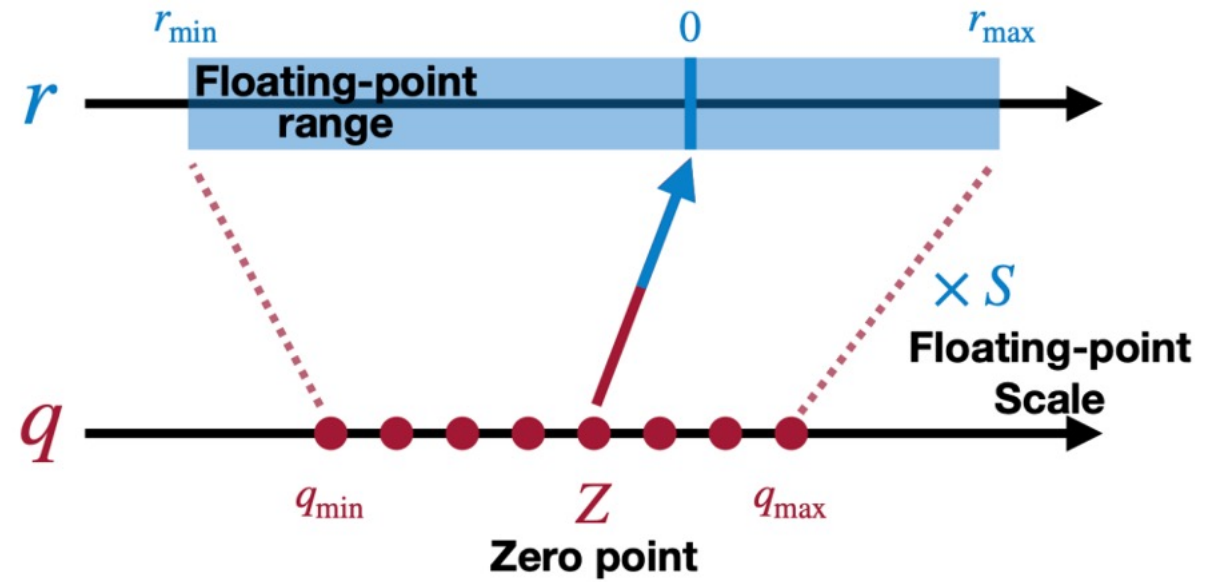
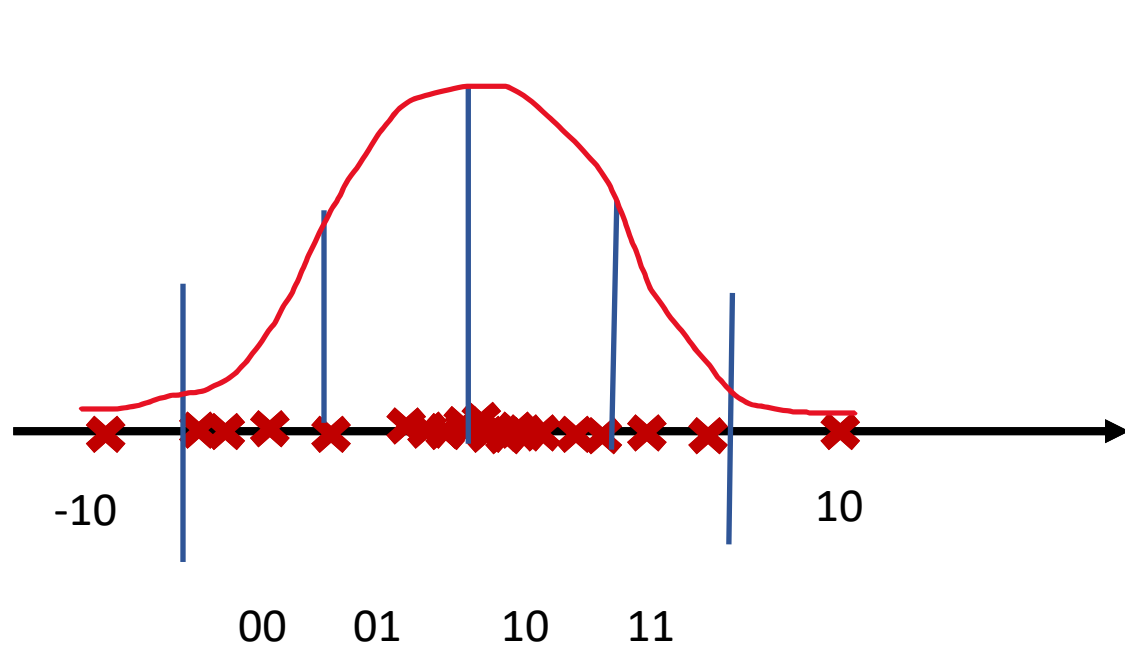
# Quantization: Data with less bits

Normal Float Quantization



Truncate the outliers!

# Quantization: Data with less bits



Binary	Decimal
01	1
00	0
11	-1
10	-2

# Quantization: Data with less bits

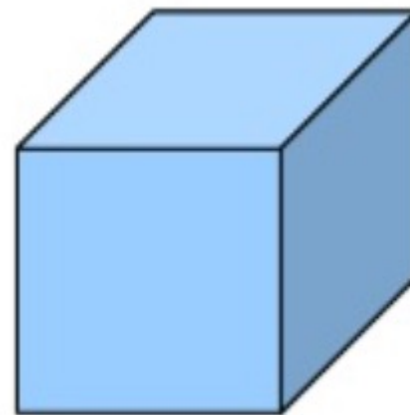
Block-wise dynamic quantization



1d-tensor



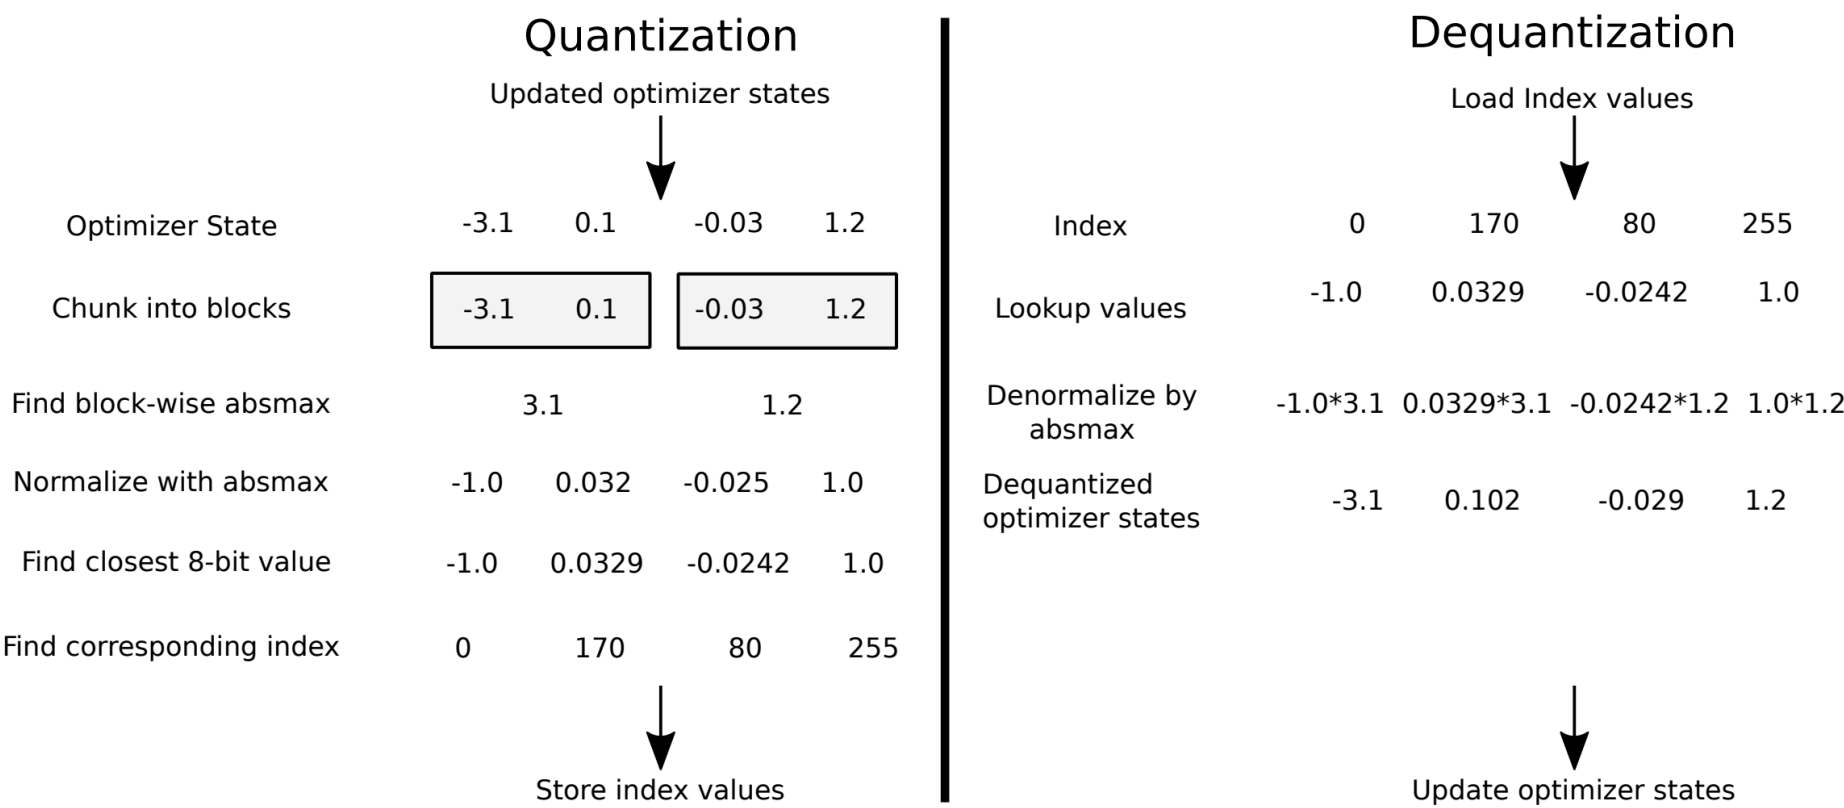
2d-tensor



3d-tensor

# Quantization: Data with less bits

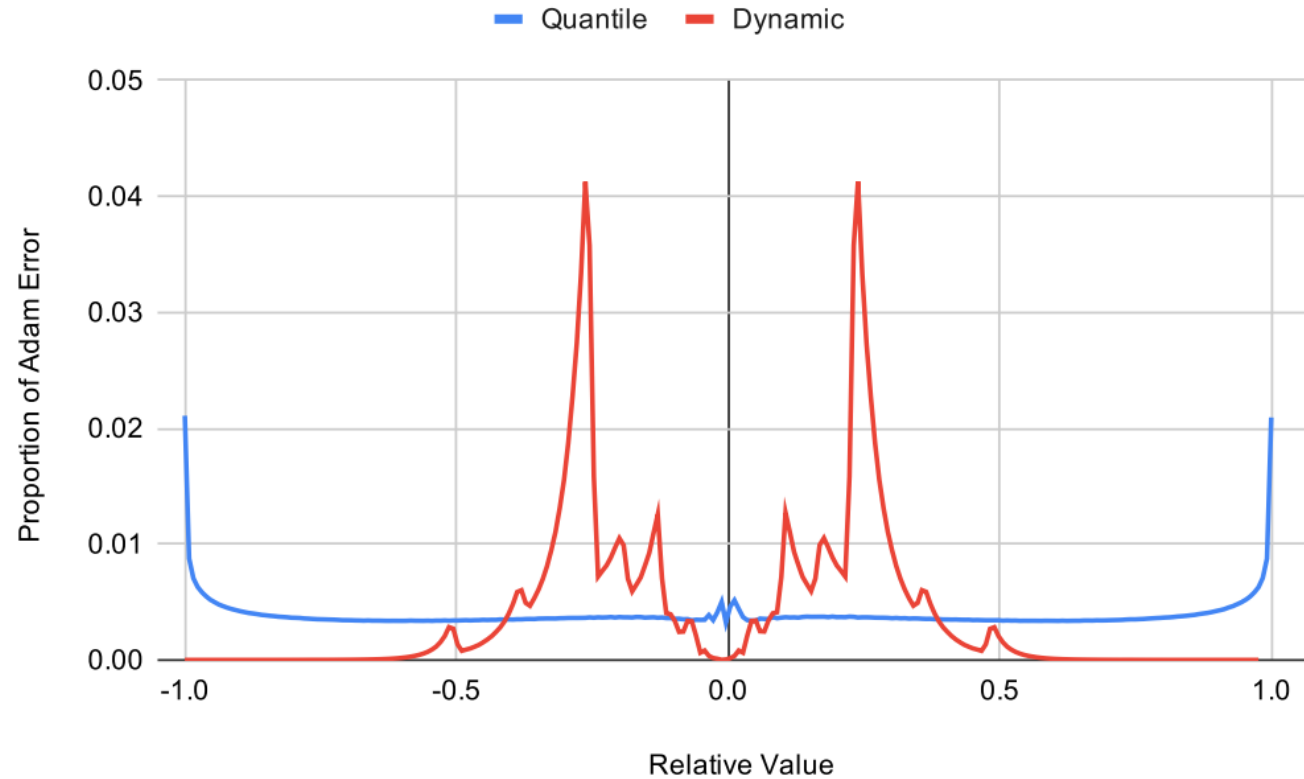
## Block-wise dynamic quantization





# Quantization: Data with less bits

Block-wise dynamic quantization



# Minimize quantization error of each layer!

Reduce the accumulated error cross the network

$$\operatorname{argmin}_{\widehat{\mathbf{W}}} ||\mathbf{W}\mathbf{X} - \widehat{\mathbf{W}}\mathbf{X}||_2^2.$$

X	2	45	-1	-17	-1
	0	12	3	-63	2
	-1	37	-1	-83	0
FP16					

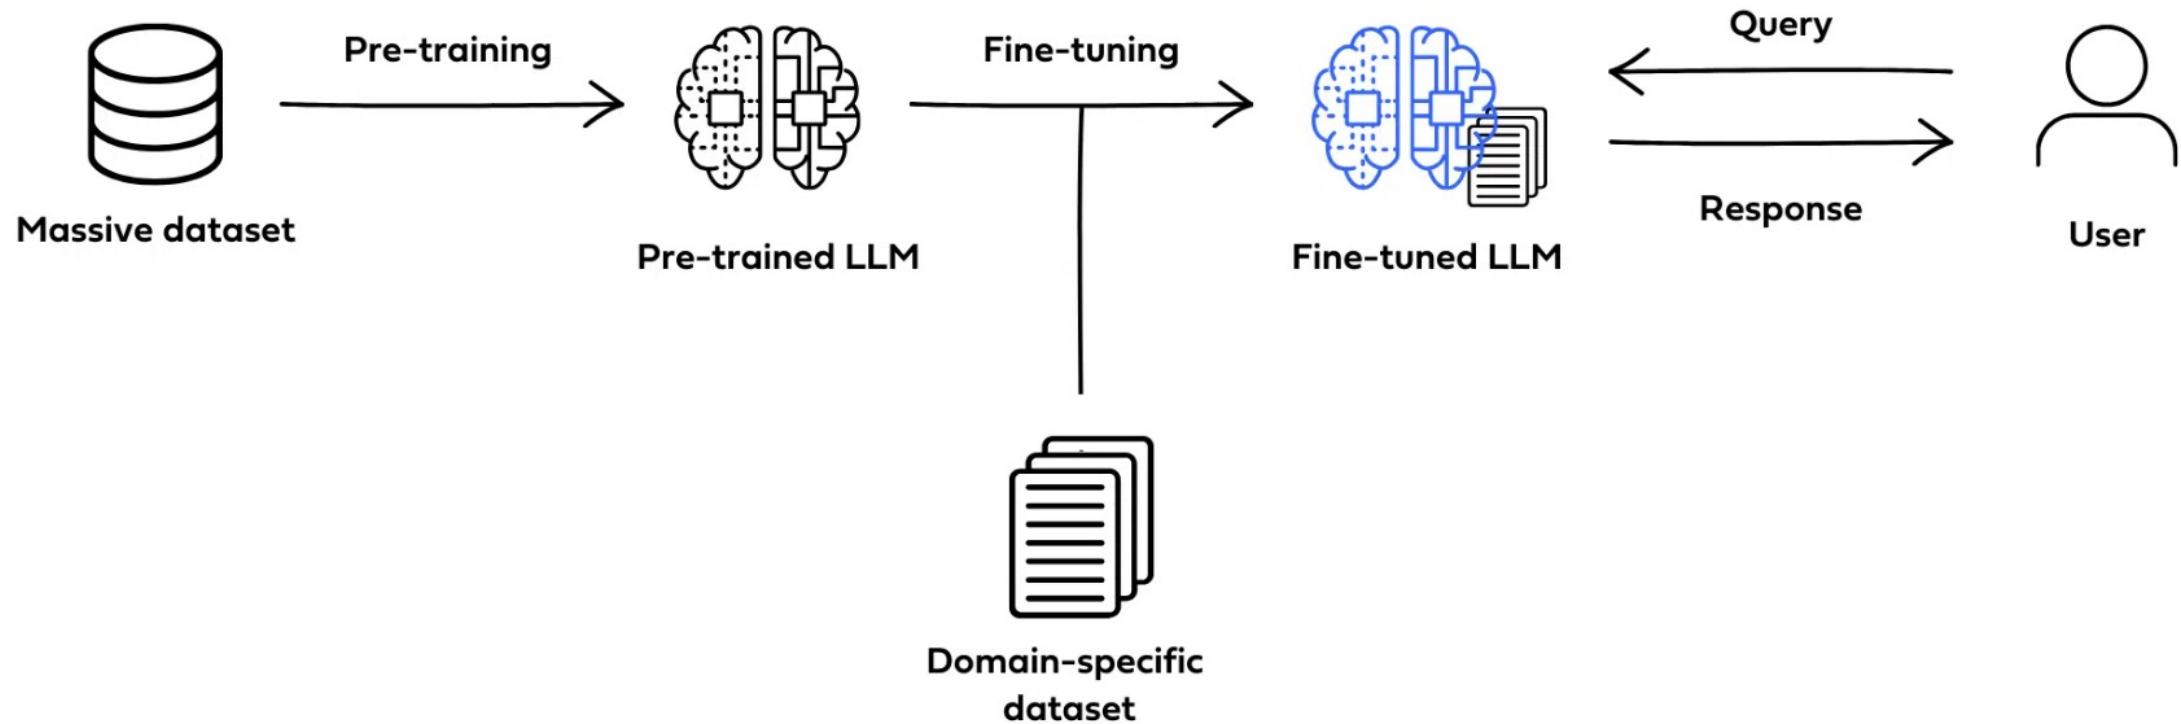
-1	0
2	0
0	-2
3	-2
-1	2
FP16	

W	-1	0
	2	0
	0	-2
3	-2	
	-1	2

# Method-1 QLORA: Efficient Finetuning of Quantized LLMs

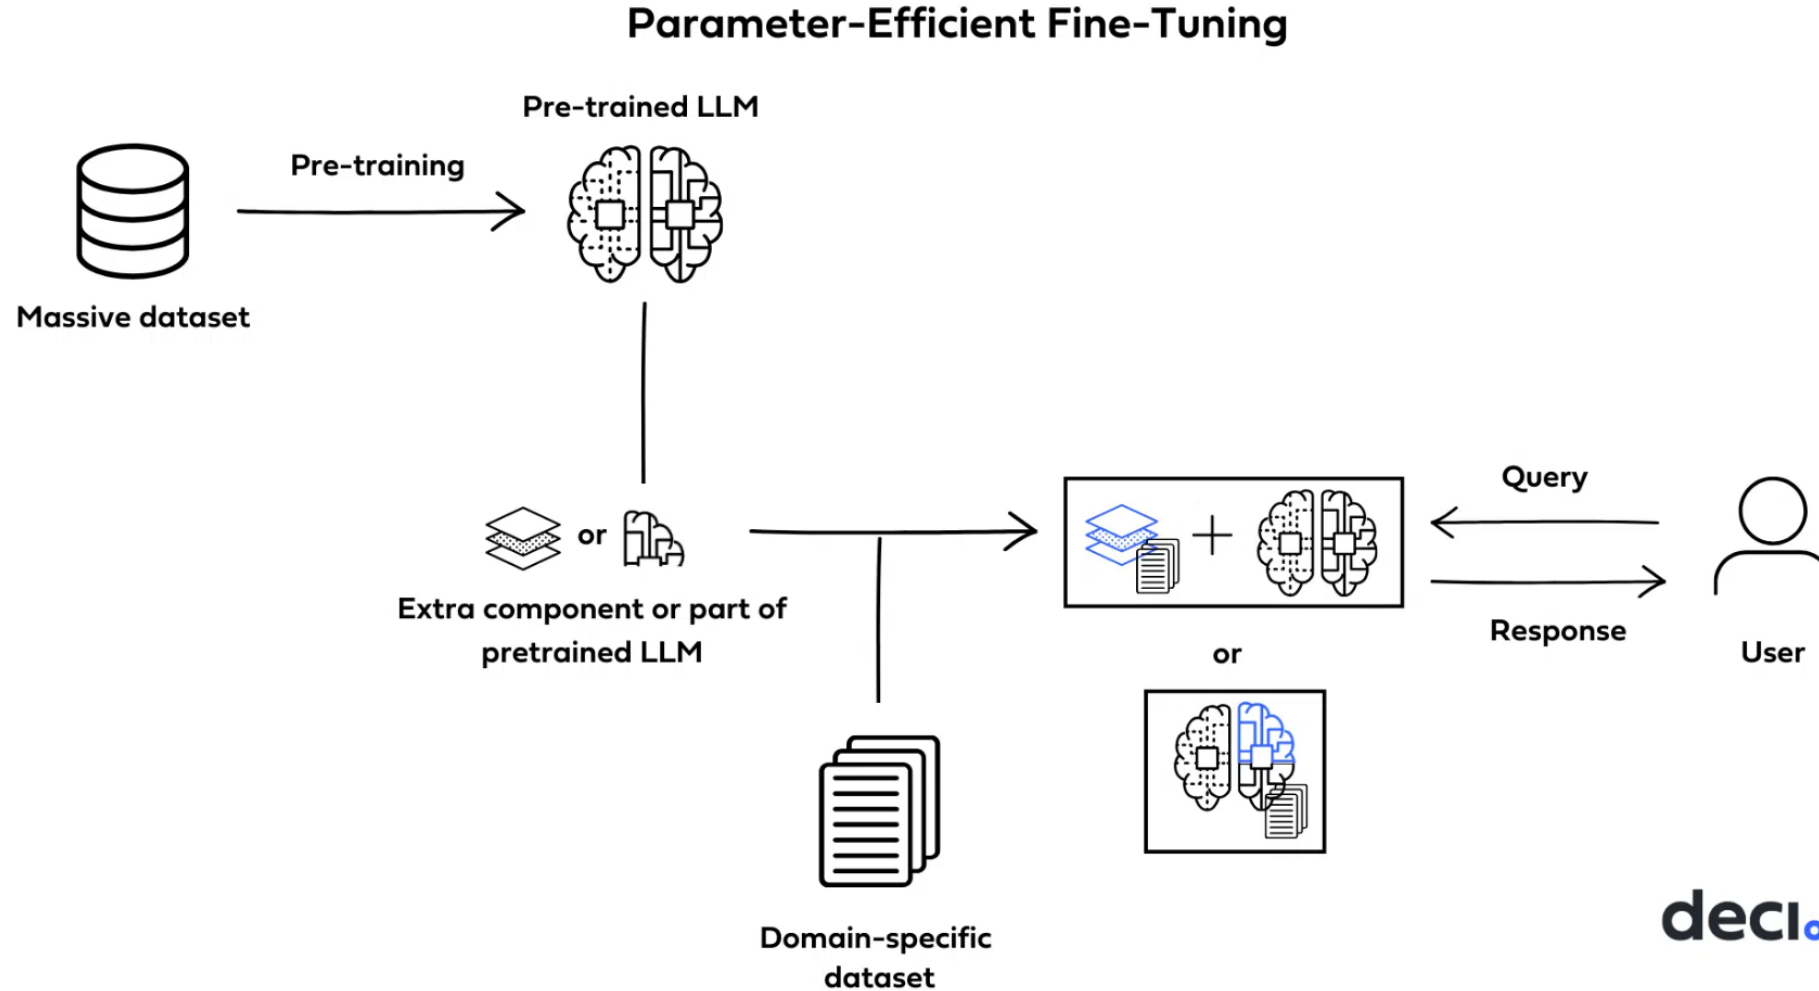
Full fine-tuning vs. Parameter-Efficient finetuning

## Full Fine Tuning



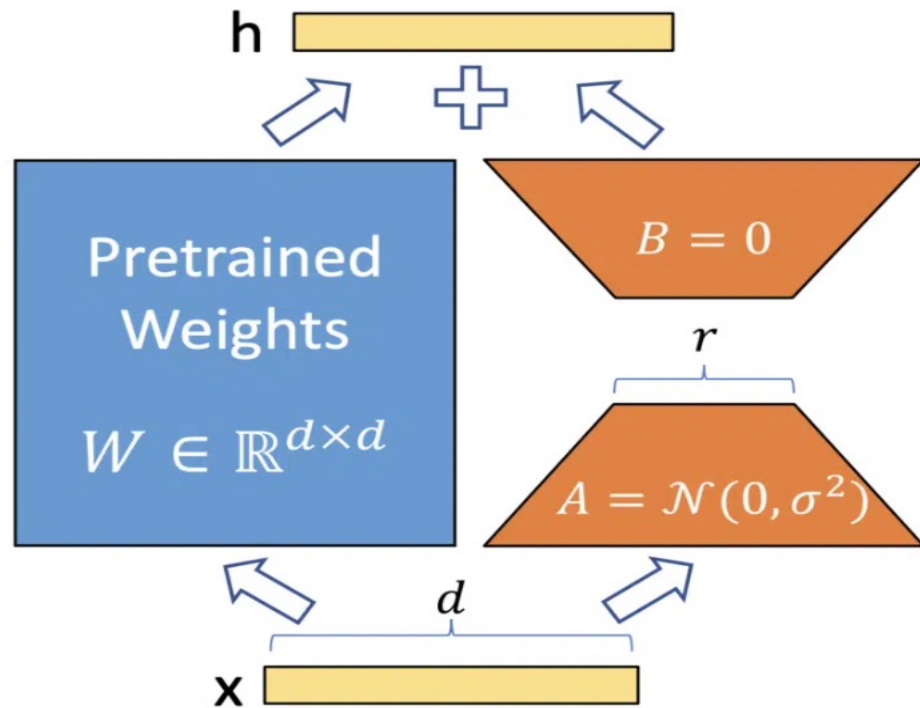
# Method-1 QLORA: Efficient Finetuning of Quantized LLMs

Full fine-tuning vs. Parameter-Efficient finetuning



## Method-1 QLORA: Efficient Finetuning of Quantized LLMs

LoRA: was introduced in early 2023, the most commonly used PEFT method

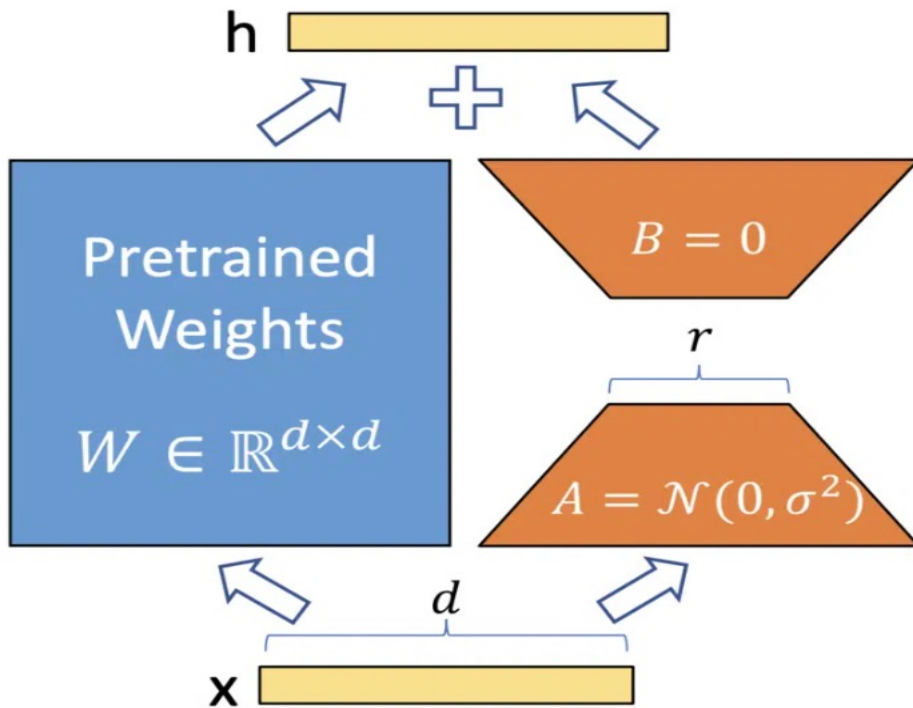


$$W' = W + A @ B$$

Figure 1: LoRA reparameterization trains only  $A$  and  $B$ . (Credit: <https://arxiv.org/abs/2106.09685>)

# Method-1 QLoRA: Efficient Finetuning of Quantized LLMs

QLoRA: the most commonly used PEFT Quantization



$$W' = \text{quant}(W) + \mathbf{A} @ \mathbf{B}$$

Figure 1: LoRA reparameterization trains only A and B. (Credit: <https://arxiv.org/abs/2106.09685>)

## *Method-1 QLORA: Efficient Finetuning of Quantized LLMs* – **Normal Float Quantization**

$$W' = \text{double\_quantize}(W) + \mathbf{A} @ \mathbf{B}$$

## Method-1 QLORA: Efficient Finetuning of Quantized LLMs – Normal Float Quantization

$$W' = \text{double\_normal\_float}(W) + A @ B$$

**QLoRA.** Using the components described above, we define QLoRA for a single linear layer in the quantized base model with a single LoRA adapter as follows:

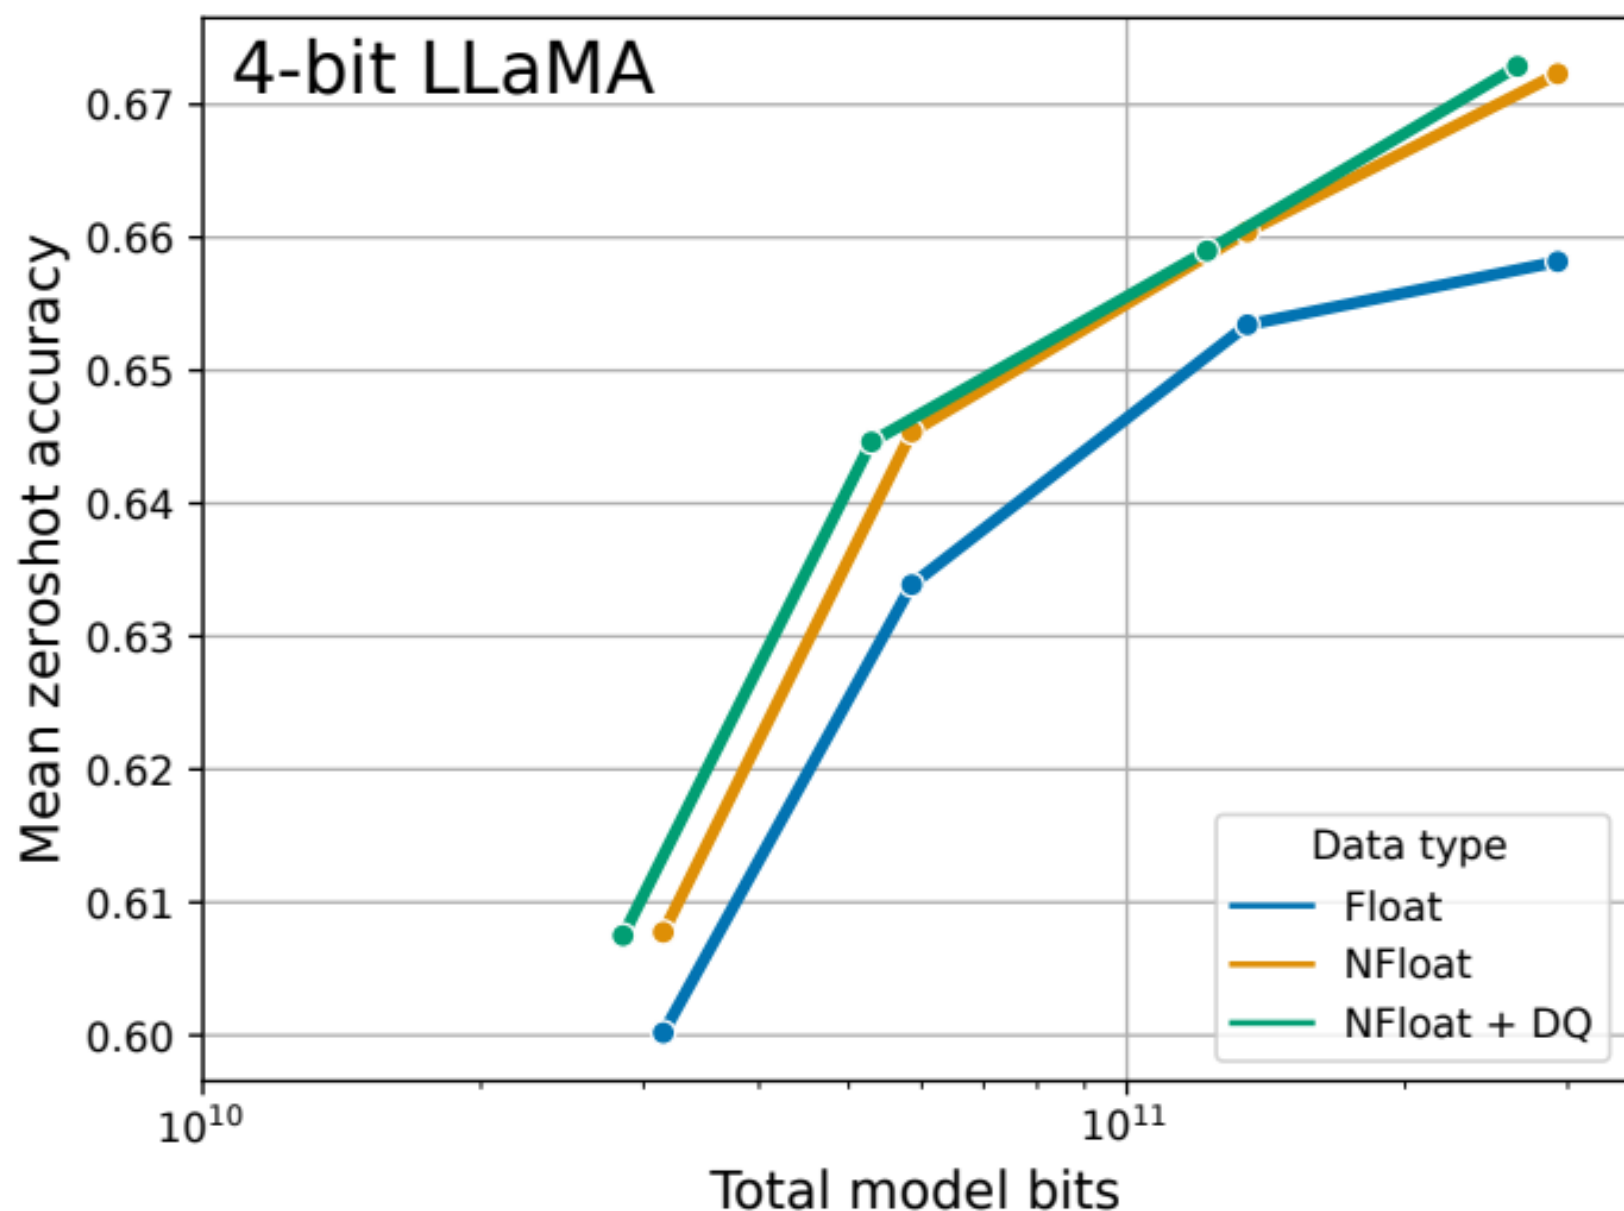
$$\mathbf{Y}^{\text{BF16}} = \mathbf{X}^{\text{BF16}} \text{doubleDequant}(c_1^{\text{FP32}}, c_2^{\text{k-bit}}, \mathbf{W}^{\text{NF4}}) + \mathbf{X}^{\text{BF16}} \mathbf{L}_1^{\text{BF16}} \mathbf{L}_2^{\text{BF16}}, \quad (5)$$

where  $\text{doubleDequant}(\cdot)$  is defined as:

$$\text{doubleDequant}(c_1^{\text{FP32}}, c_2^{\text{k-bit}}, \mathbf{W}^{\text{k-bit}}) = \text{dequant}(\text{dequant}(c_1^{\text{FP32}}, c_2^{\text{k-bit}}), \mathbf{W}^{\text{4bit}}) = \mathbf{W}^{\text{BF16}}, \quad (6)$$

We use NF4 for  $\mathbf{W}$  and FP8 for  $c_2$ . We use a blocksize of 64 for  $\mathbf{W}$  for higher quantization precision and a blocksize of 256 for  $c_2$  to conserve memory.



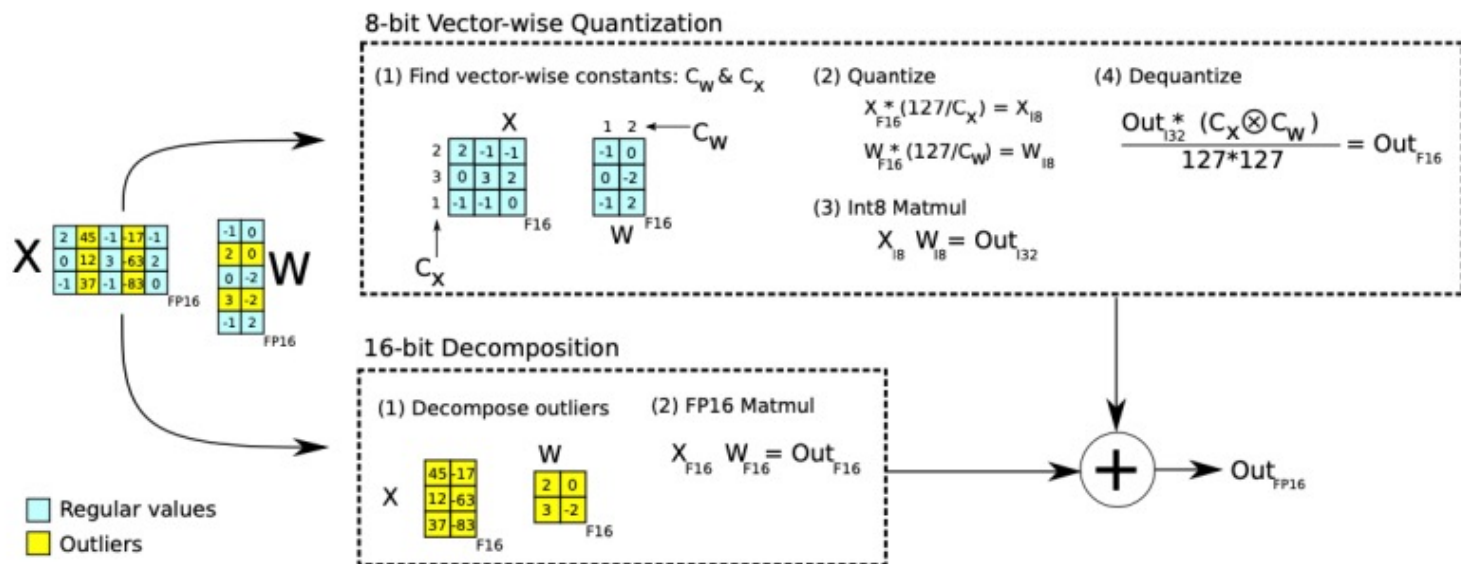
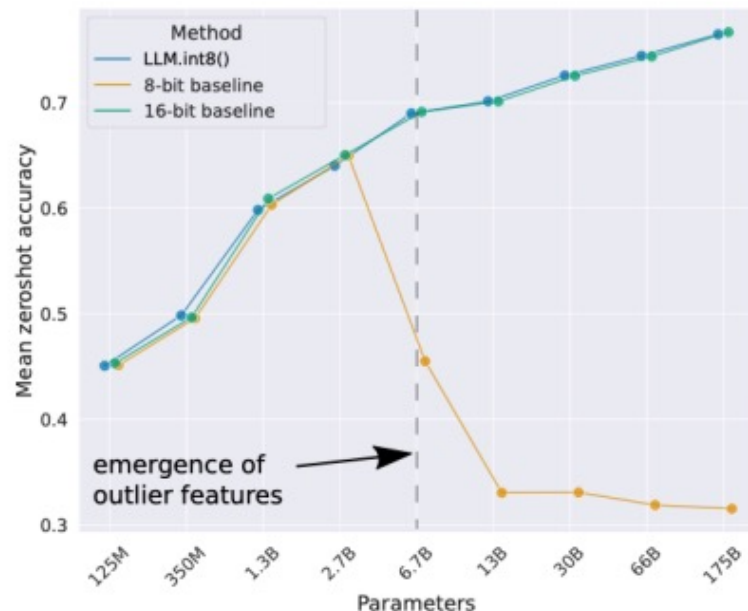


**Table 3:** Experiments comparing 16-bit BrainFloat (BF16), 8-bit Integer (Int8), 4-bit Float (FP4), and 4-bit NormalFloat (NF4) on GLUE and Super-NaturalInstructions. QLoRA replicates 16-bit LoRA and full-finetuning.

Dataset Model	GLUE (Acc.)	Super-NaturalInstructions (RougeL)				
	RoBERTa-large	T5-80M	T5-250M	T5-780M	T5-3B	T5-11B
BF16	88.6	40.1	42.1	48.0	54.3	62.0
BF16 replication	88.6	40.0	42.2	47.3	54.9	-
LoRA BF16	88.8	40.5	42.6	47.1	55.4	60.7
QLoRA Int8	88.8	40.4	42.9	45.4	56.5	60.7
QLoRA FP4	88.6	40.3	42.4	47.5	55.6	60.9
QLoRA NF4 + DQ	-	40.4	42.7	47.7	55.3	60.9

# Quantization — LLM.int8()

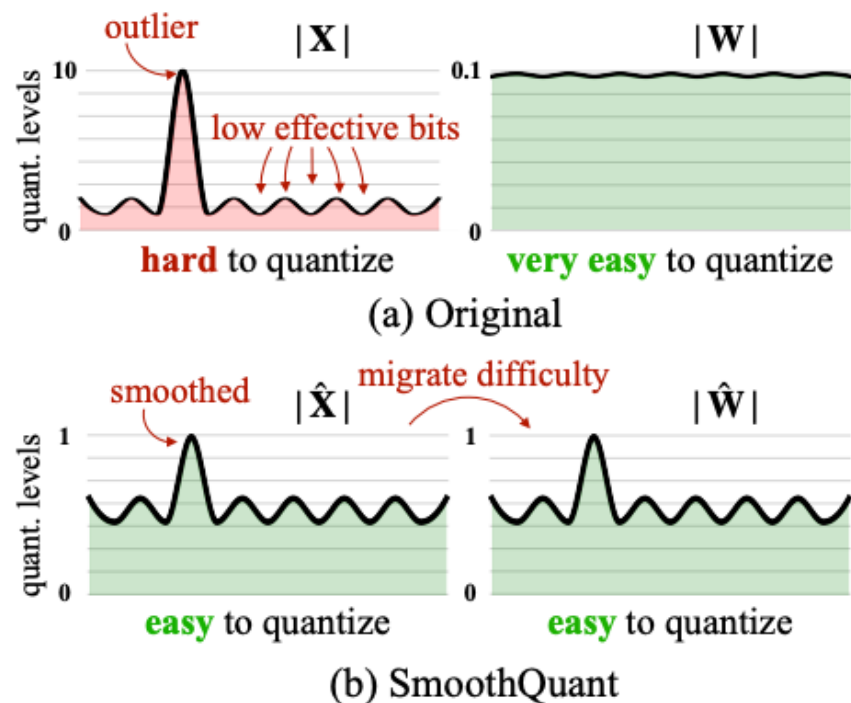
## Mixed-Precision Decomposition



- **Motivation:** Transformers have outlier features that have **large values** (especially large models).
  - They occur in particular hidden dimensions, leading to large quantization error.
- **Key idea:** Separate outlier features into a **separate FP16 MM**, quantize the other values to Int8.
  - Outlier: At least one feature dimension with a magnitude larger than the threshold (6).
  - Token-wise scale factor (for  $X$ ) and (output) channel-wise scale factor (for  $W$ ).

# SmoothQuant: Accurate and Efficient Post-Training Quantization for LLMs

	LLM (100B+) Accuracy	Hardware Efficiency
ZeroQuant	✗	✓
Outlier Suppression	✗	✓
LLM.int8()	✓	✗
<b>SmoothQuant</b>	✓	✓



- We propose SmoothQuant, an **accurate** and **efficient** post-training-quantization (PTQ) method to enable 8-bit weight, 8-bit activation (**W8A8**) quantization for LLMs.
- Since **weights are easy** to quantize while **activations are not**, SmoothQuant smooths the activation outliers by **migrating the quantization difficulty from activations to weights** with a mathematically equivalent transformation.



